

Image Classification

Notes based on

Machine Learning for the Web by YiNing at ITP, NYU,

How to build a Teachable Machine with TensorFlow.js by Nikhil Thorat from Google Brain, PAIR,

and Daniel Shiffman from the Coding Train

Classification: Basic Human task



Classification: Basic Human task

Input: image



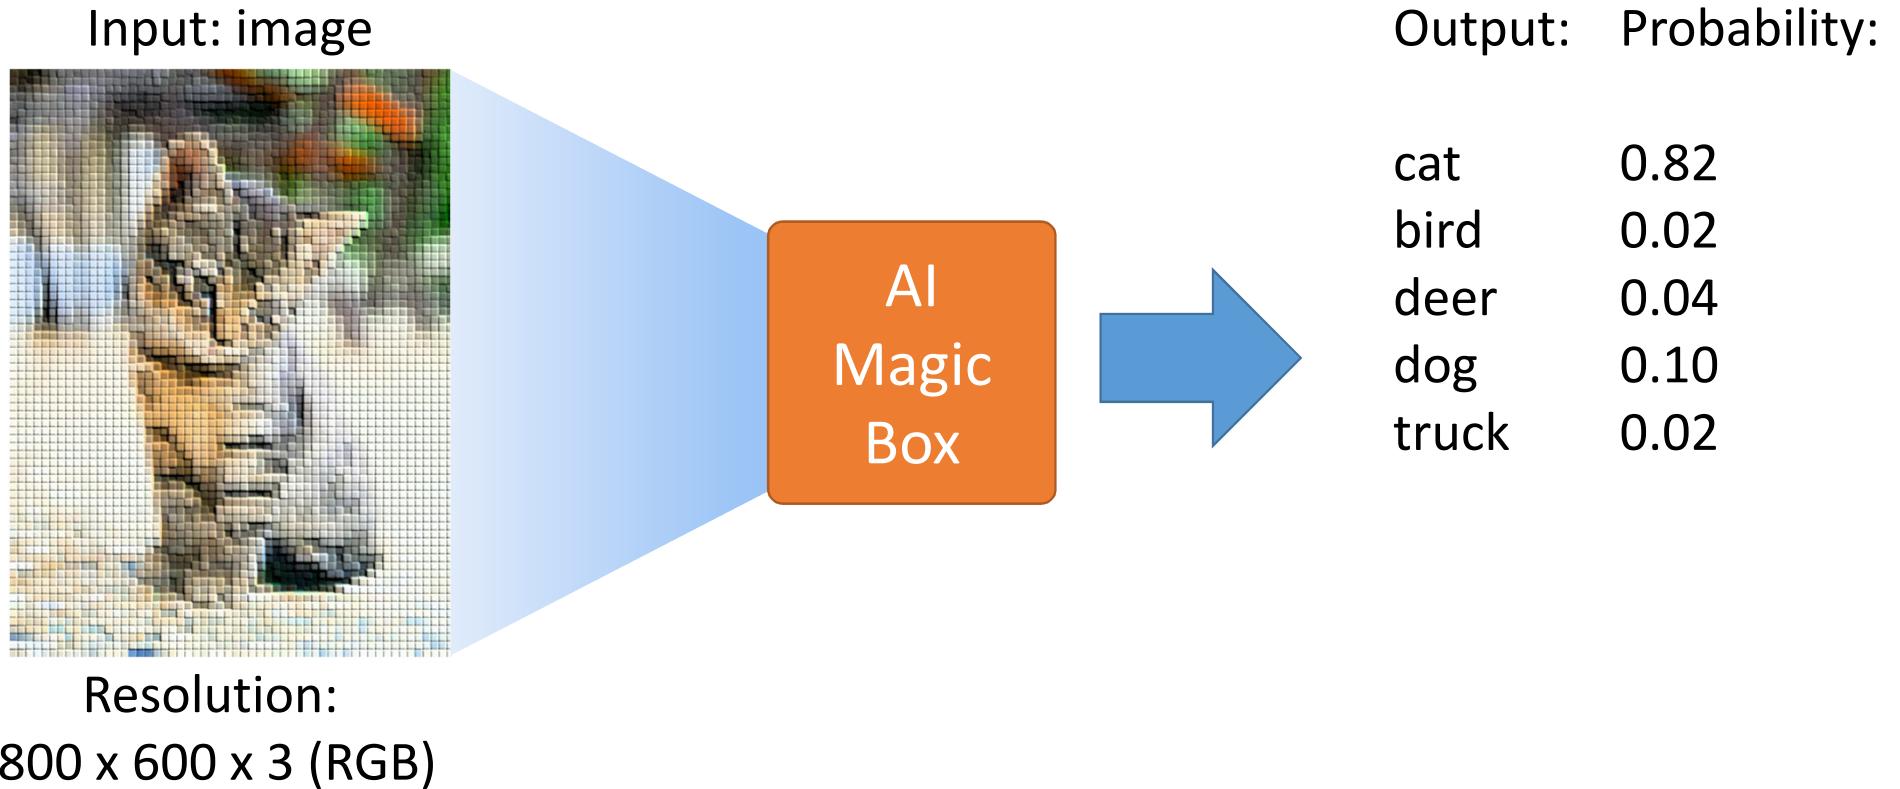
Output:

cat
bird
deer
dog
truck

Resolution:
800 x 600 x 3 (RGB)

All images from <https://unsplash.com/s/photos/cat>

Classification: Basic AI task



AI Classification: Challenges

- Change view angle



AI Classification: Challenges

- Variation within class



AI Classification: Challenges

- Sub-classes (different cat breeds)



Maine Coon



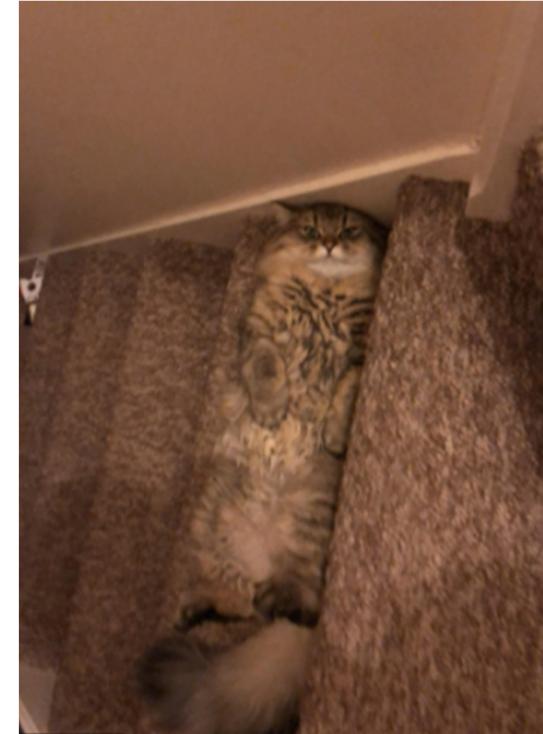
British Shorthair



Grumpy Cat
(Mixed breed)
2012-2019

AI Classification: Challenges

- Background blending / camouflage



AI Classification: Challenges

- Lighting change



AI Classification: Challenges

- Deformation



AI Classification: Challenges

- Occlusion



Image Classification: Building Blocks for other tasks.

- Object detection



Output: bounding boxes and class

Image Classification: Building Blocks for other tasks.

- Object segmentation



Output: segments and class

Image Classification: Building Blocks for other tasks.

- Image captioning



Output: description: There are two cars and a bike on the road

Image Classification: Building Blocks for other tasks.

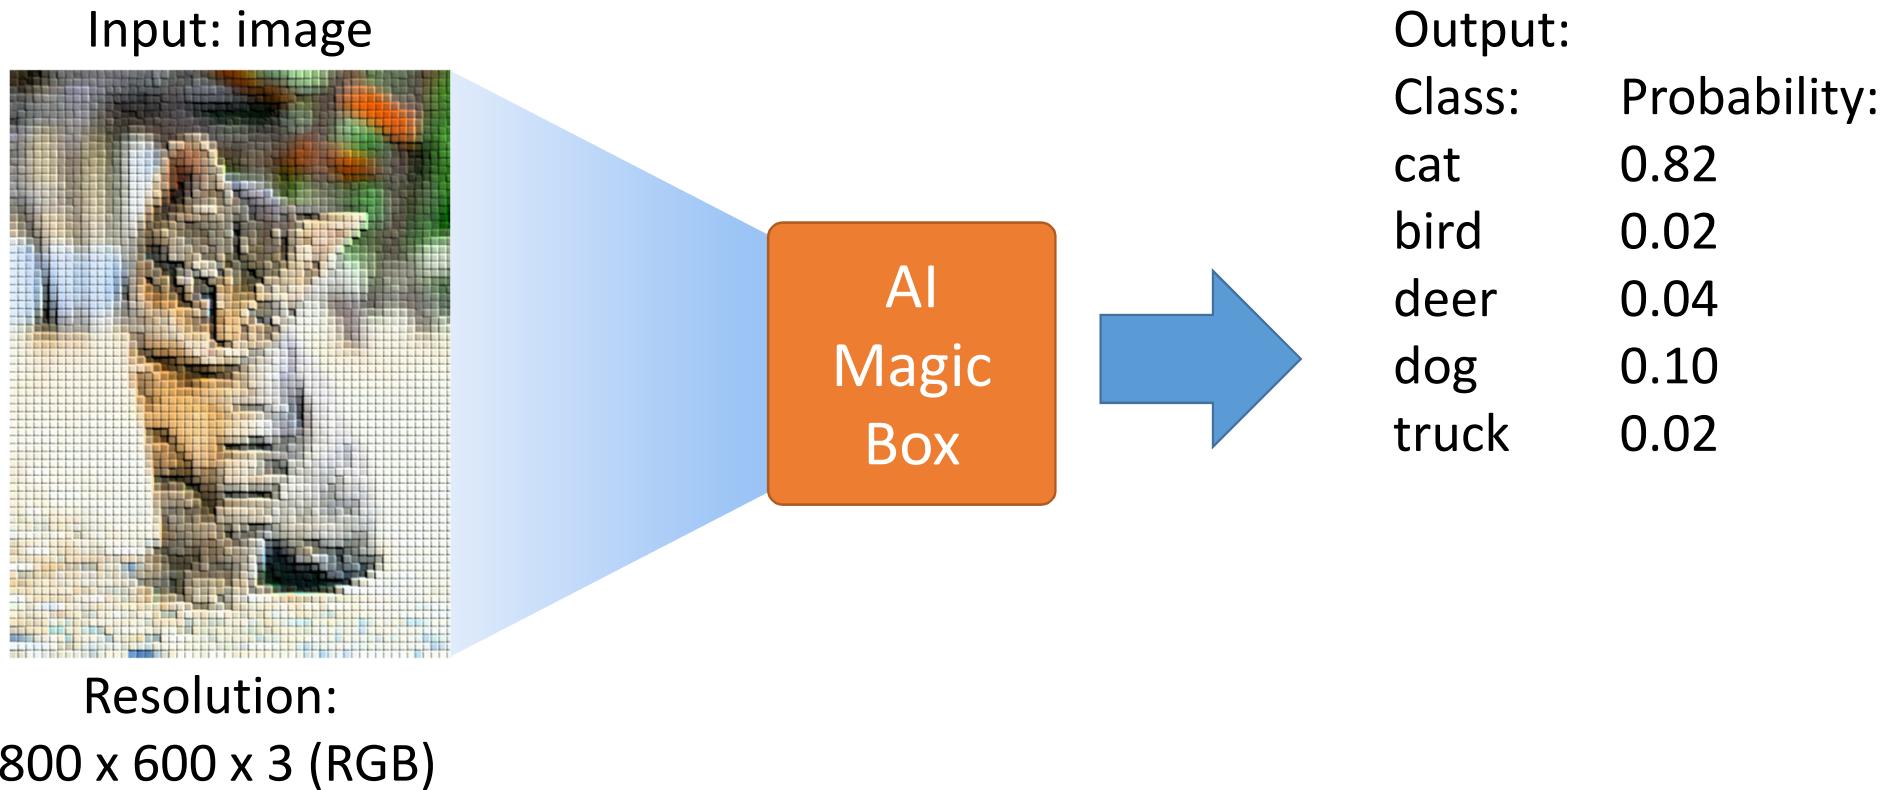
- Playing go



(1, 1)
(1, 2)
...
(1, 19)
...
(19, 19)

What's the
next move?

Classification: Basic AI task



All images from <https://unsplash.com/s/photos/cat>

An Image Classifier

```
def classify_image(image):  
    # some AI magic  
    return class_label
```

Unlike e.g. sorting a list of numbers,

no obvious way to hard-code the algorithm for
recognizing a cat, or other classes.

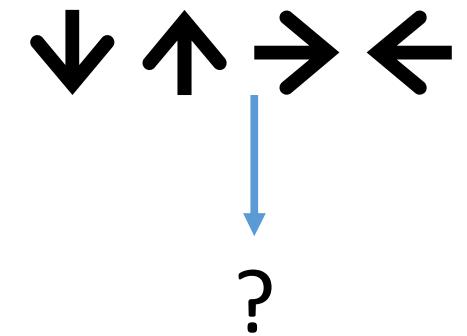
You can try edge detection...



Find edges



Find corners



Machine Learning: Data driven approach

1. Dataset – Collect a dataset of images with labels
2. Train – Use ML to train classifier
3. Test/deploy – evaluate classifier on new images

```
def train(images, labels):  
    # Machine learning!  
    return model  
  
def predict(model, test_image):  
    # use model to predict label  
    return predicted_label
```

Example training set

airplane



automobile



bird



cat



deer



Image Classification Datasets: MNIST



10 classes: Digits 0 to 9

28x28 grayscale images

50k training images

10k test images

Results from MNIST often do not hold on more complex datasets!

Image Classification Datasets: Fashion MNIST

| Label | Description | Examples |
|-------|-------------|----------|
| 0 | T-Shirt/Top | |
| 1 | Trouser | |
| 2 | Pullover | |
| 3 | Dress | |
| 4 | Coat | |
| 5 | Sandals | |
| 6 | Shirt | |
| 7 | Sneaker | |
| 8 | Bag | |
| 9 | Ankle boots | |

10 classes

28x28 grayscale images

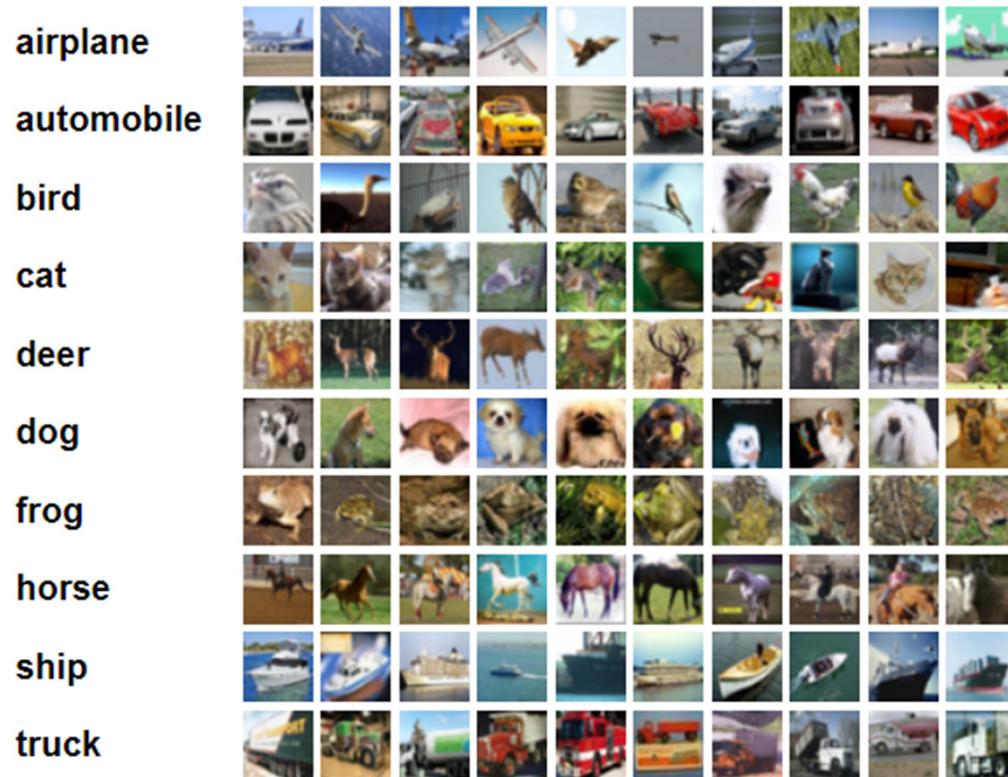
50k training images

10k test images

<https://github.com/zalandoresearch/fashion-mnist>

Results from MNIST often do not hold on more complex datasets!

Image Classification Datasets: CIFAR10



10 classes

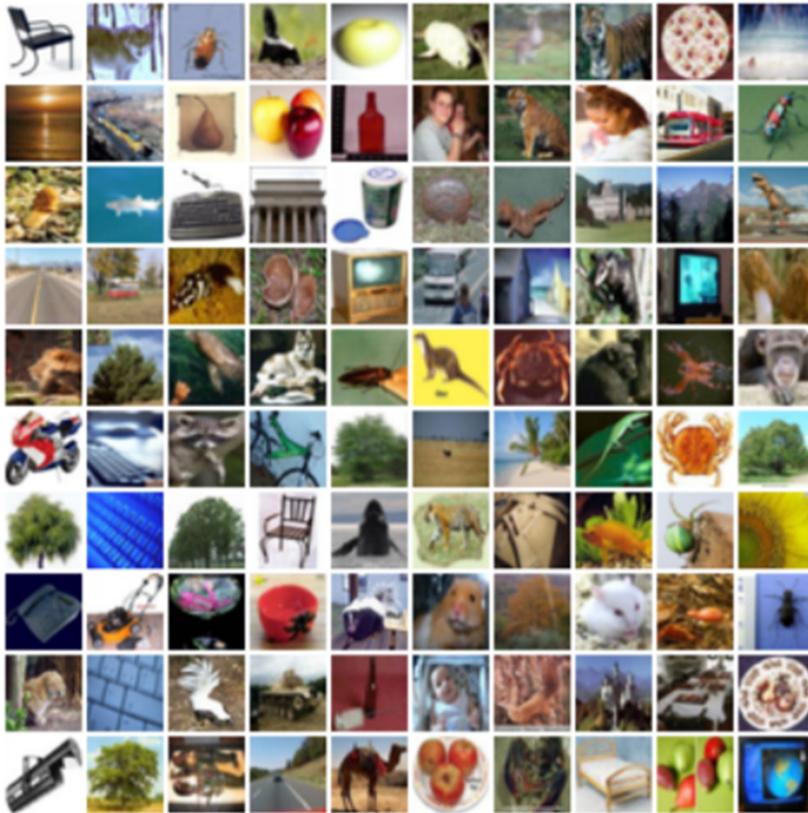
50k training images (5k per class)

10k testing images (1k per class)

32x32 RGB images

<https://www.cs.toronto.edu/~kriz/cifar.html>

Image Classification Datasets: CIFAR100



<https://www.cs.toronto.edu/~kriz/cifar.html>

100 classes

50k training images (500 per class)

10k testing images (100 per class)

32x32 RGB images

20 **superclasses** with 5 classes each:

Aquatic mammals:

beaver, dolphin, otter, seal, whale

Trees:

Maple, oak, palm, pine, willow

Image Classification Datasets: ImageNet



1000 classes

~1.3M training images (~1.3K per class)
50K validation images (50 per class)
100K test images (100 per class)

Performance metric: **Top 5 accuracy**

Algorithm predicts 5 labels for each image; one of them needs to be right

Image Classification Datasets: MIT Places



365 classes of different scene types

~8M training images

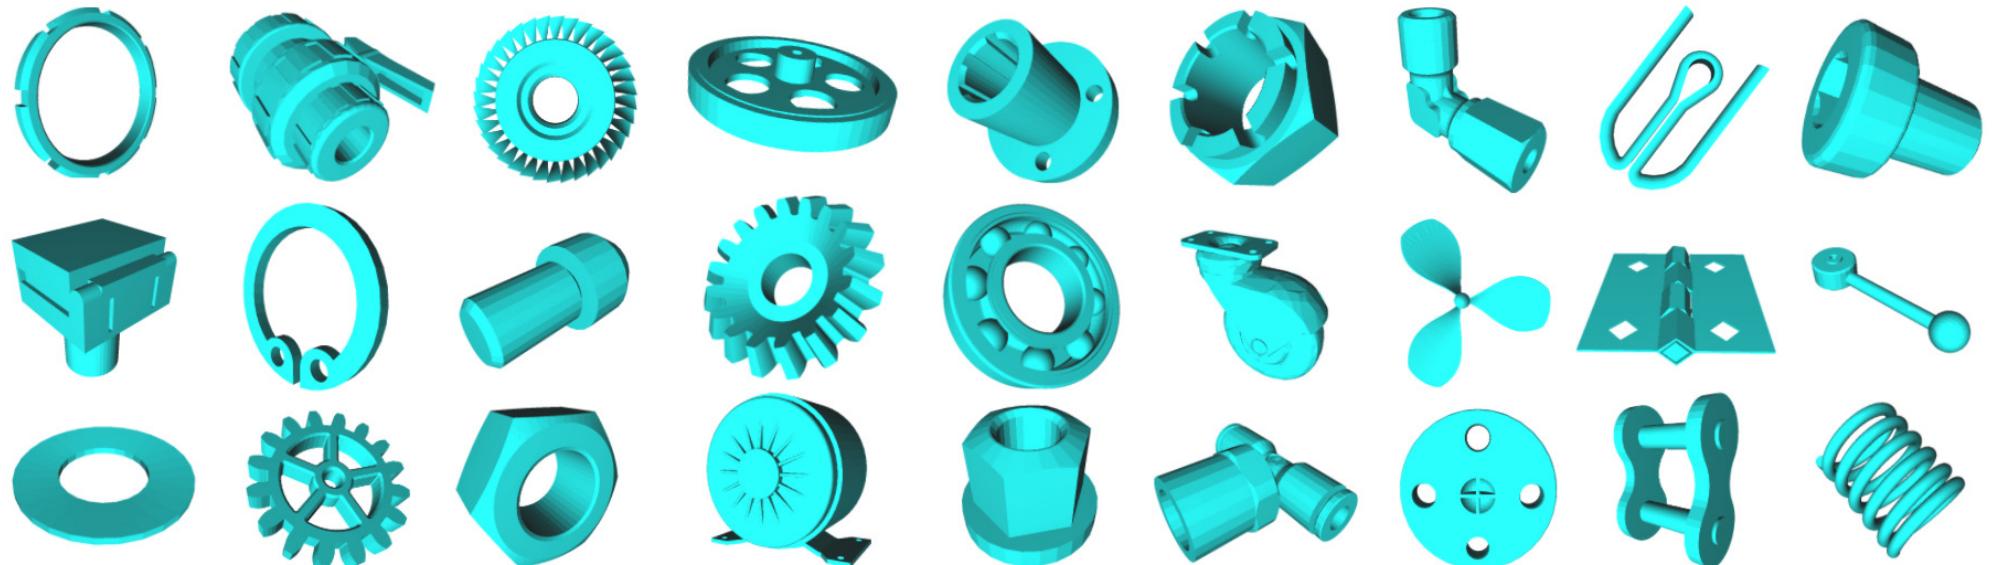
18.25K val images (50 per class)

328.5K test images (900 per class)

Images have variable size, often resized to **256x256** for training

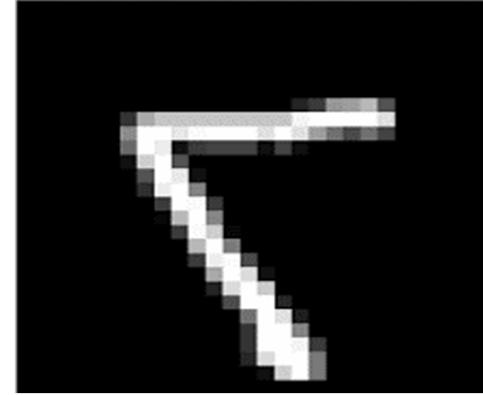
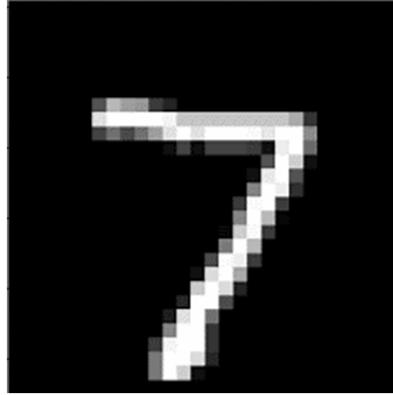
A Large-scale Mechanical Components Benchmark for Deep Neural Networks

- <https://mechanical-components.herokuapp.com/>
- <https://www.youtube.com/watch?v=EaKo7ky15uA>

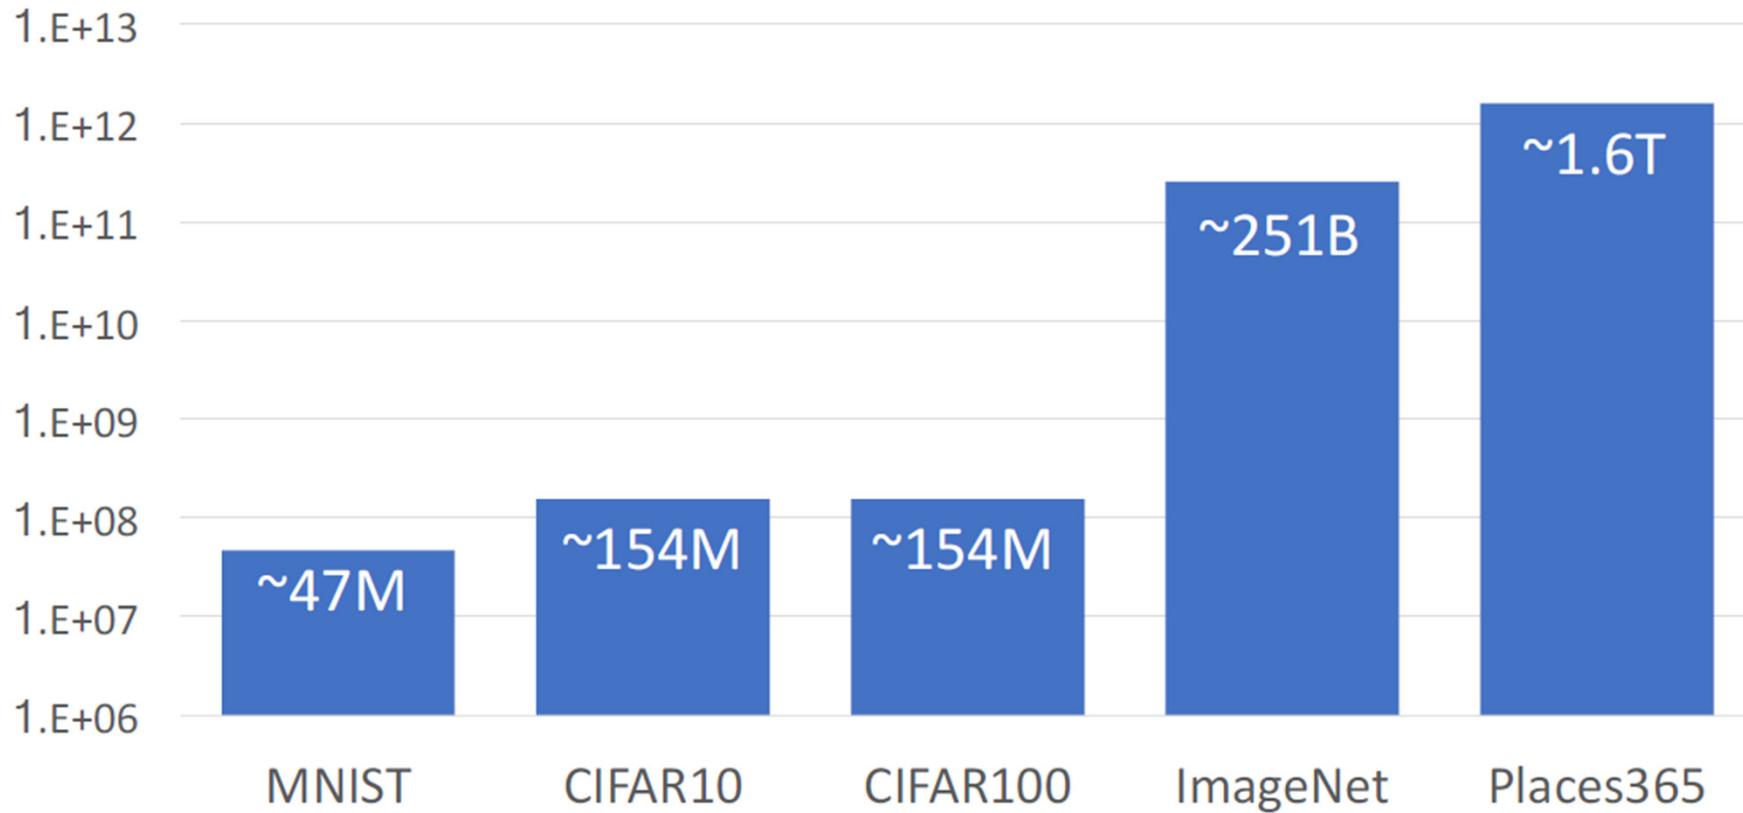


Datasets

1. MNIST
2. CIFAR10
3. CIFAR100
4. ImageNet
5. MIT Places



Datasets: Number of Training Pixels



Your first classifier: Nearest Neighbour

```
def train(images, labels):  
    # Machine learning!  
    return model
```



Memorize all data
and labels

```
def predict(model, test_image):  
    # use model to predict label  
    return predicted_label
```



Predict the label of
the most similar
training image

Distance Metric to compare images

L1 distance:

$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$

test image

| | | | |
|----|----|-----|-----|
| 56 | 32 | 10 | 18 |
| 90 | 23 | 128 | 133 |
| 24 | 26 | 178 | 200 |
| 2 | 0 | 255 | 220 |

training image

| | | | |
|----|----|-----|-----|
| 10 | 20 | 24 | 17 |
| 8 | 10 | 89 | 100 |
| 12 | 16 | 178 | 170 |
| 4 | 32 | 233 | 112 |

-

pixel-wise absolute value differences

| | | | |
|----|----|----|-----|
| 46 | 12 | 14 | 1 |
| 82 | 13 | 39 | 33 |
| 12 | 10 | 0 | 30 |
| 2 | 32 | 22 | 108 |

=

add
→ 456

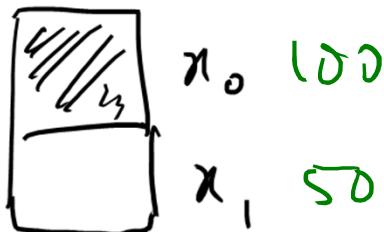
What does this look like?



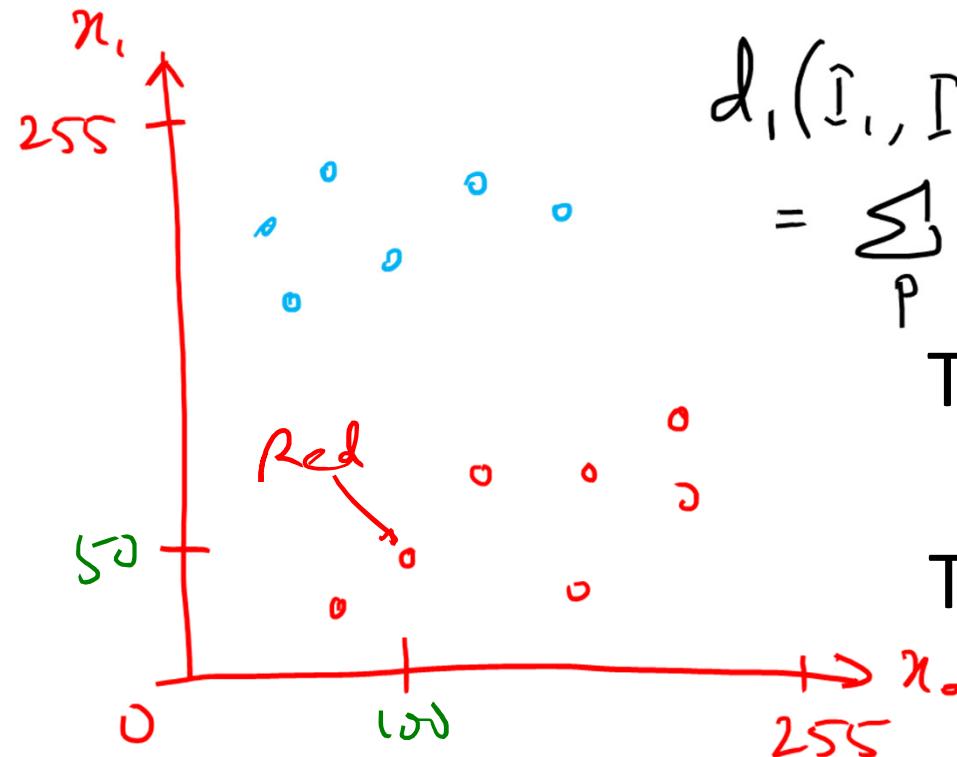
How to find the nearest neighbour?

- Distance

Consider a 2-pixel image



pixel values : 0 - 255



$$d_1(\hat{I}_1, \hat{I}_2)$$

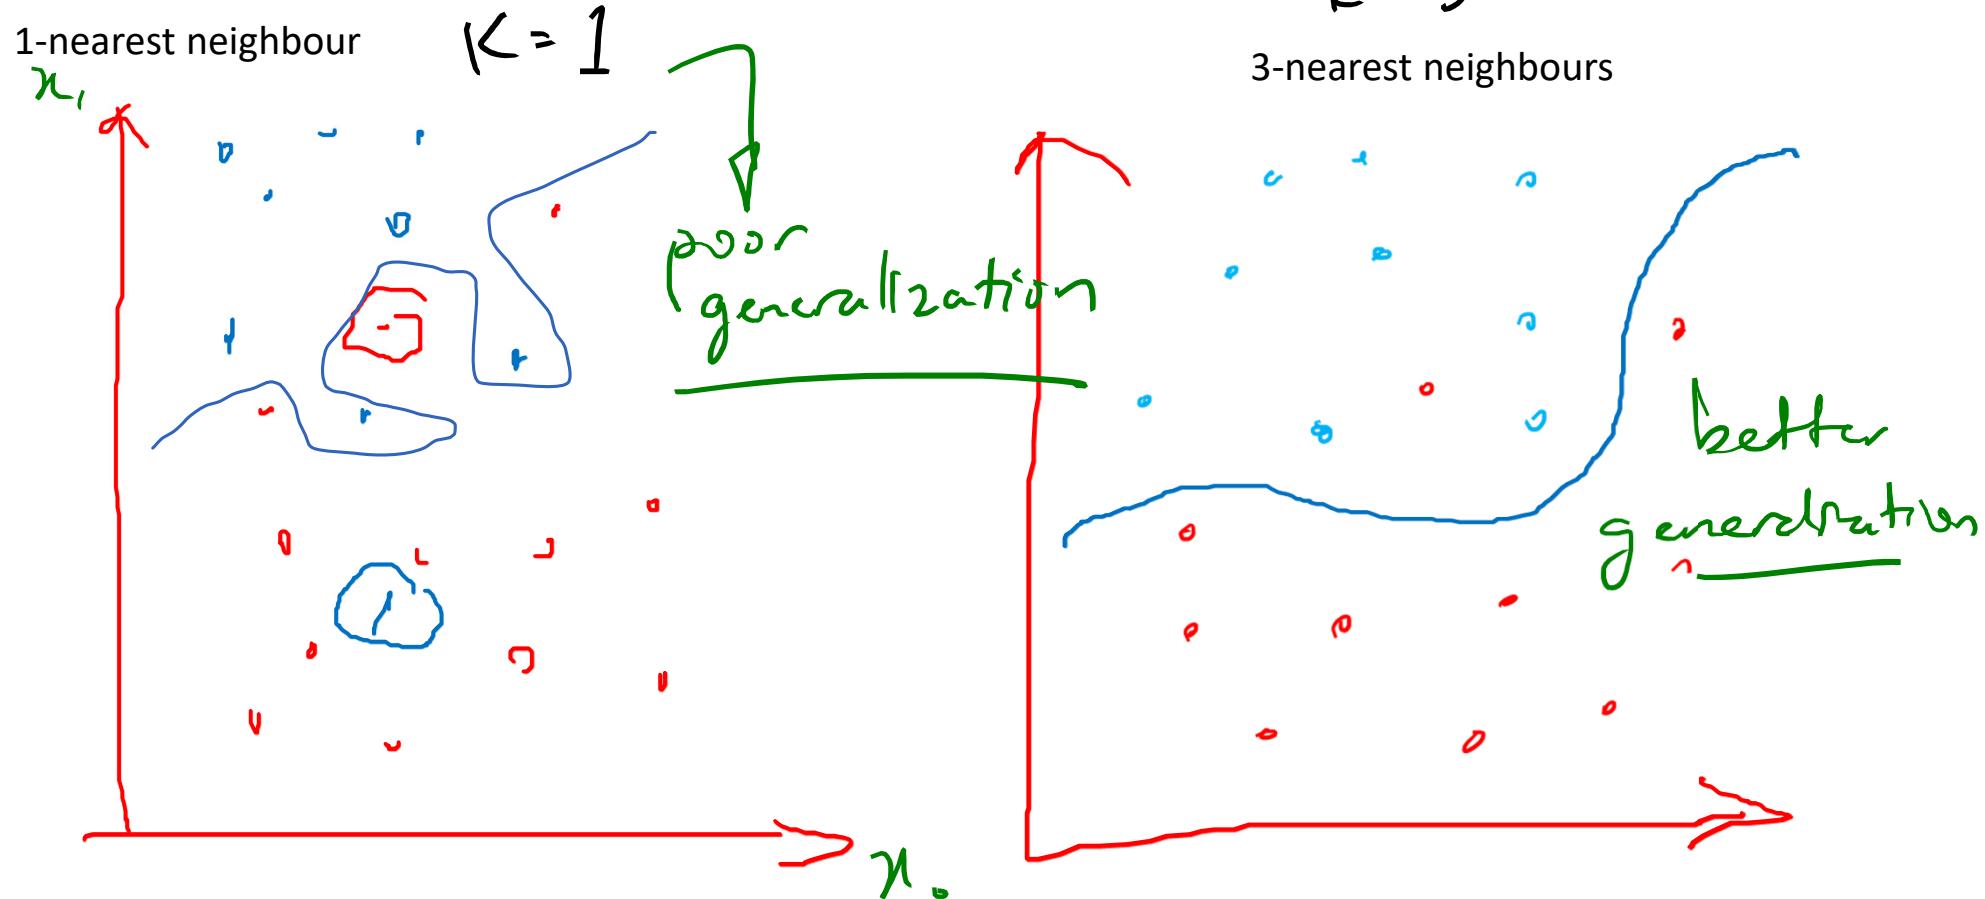
$$= \sum_p |I_1^p - I_2^p|$$

Training Time = N

vs

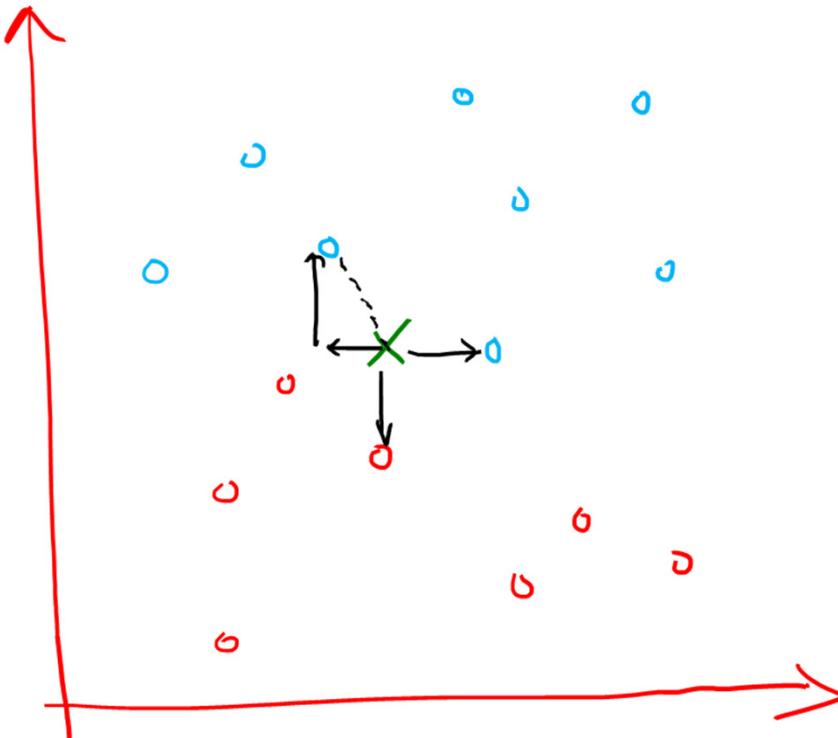
Testing Time = N²

Nearest neighbour Decision Boundaries



K-Nearest Neighbour

$$k = 3$$



K-Nearest Neighbour – Distance Calculation

- Manhattan
- Euclidean

$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$

$$d_2(I_1, I_2) = \sqrt{\sum_p (I_1^p - I_2^p)^2}$$

Hyperparameters

- What is the best value of **K** to use?
- What is the best **distance metric** to use?

These are examples of **hyperparameters**: choices about our learning algorithm that we don't learn from the training data; instead we set them at the start of the learning process

Very problem-dependent. In general need to try them all and see what works best for our data / task.

Setting Hyperparameters

Idea #1: Choose hyperparameters that work best on the data

BAD: $K = 1$ always works perfectly on training data

Your dataset

Setting Hyperparameters

Idea #1: Choose hyperparameters that work best on the data

BAD: K = 1 always works perfectly on training data

Your dataset

Idea #2: Split data into **train** and **test**, choose hyperparameters that work best on test data

BAD: No idea how algorithm will perform on new data

train

test

Setting Hyperparameters

Idea #1: Choose hyperparameters that work best on the data

BAD: K = 1 always works perfectly on training data

Your dataset

Idea #2: Split data into **train** and **test**, choose hyperparameters that work best on test data

BAD: No idea how algorithm will perform on new data

train

test

Idea #3: Split data into **train**, **val**, and **test**; choose hyperparameters on val and evaluate on test

Better!

train

validate

test

K-Nearest Neighbour

- Seldom used at test time
- Distance metrics on pixels are not informative
- However, K-Nearest Neighbour combined with ConvNet (feature extractor) works very well

MobileNet



- Use ml5.js on p5.js web editor
- Loaded MobileNet model
- Classified the input image:
cock

What if we want to define our own classes?

Projects related to Teachable Machine

- [Pacman](#)
- [Pong game](#)
- [Objectifier Spatial Programming](#)
- [Project Euphonia](#)
- [Teachable Snake](#)
- [Teachable Sorter](#)
- [Teachable Arcade](#)
- [Eyeo Festival 2019 - Coding train, code](#)

How does teachable machine work?

Image from Webcam → [MobileNet](#)(Second to last / last layer) → [KNN Classifier](#) → Output

High-level semantic features of the image

MobileNet Model

- **MobileNet: Efficient Convolutional Neural Networks for Mobile Vision Applications**([paper](#)).



Model Loaded

The MobileNet model labeled this as robin, American robin, *Turdus migratorius*, with a confidence of 0.99.



| class name | probability | imagenet class id |
|------------------------|-------------|-------------------|
| Egyptian cat | 0.478 | 285 |
| tabby, tabby cat | 0.300 | 281 |
| tiger cat | 0.167 | 282 |
| remote control, remote | 0.016 | 761 |
| Siamese cat, Siamese | 0.008 | 284 |

MobileNet Model

- **Small, fast, accurate.** A model that is trained on **ImageNet**.

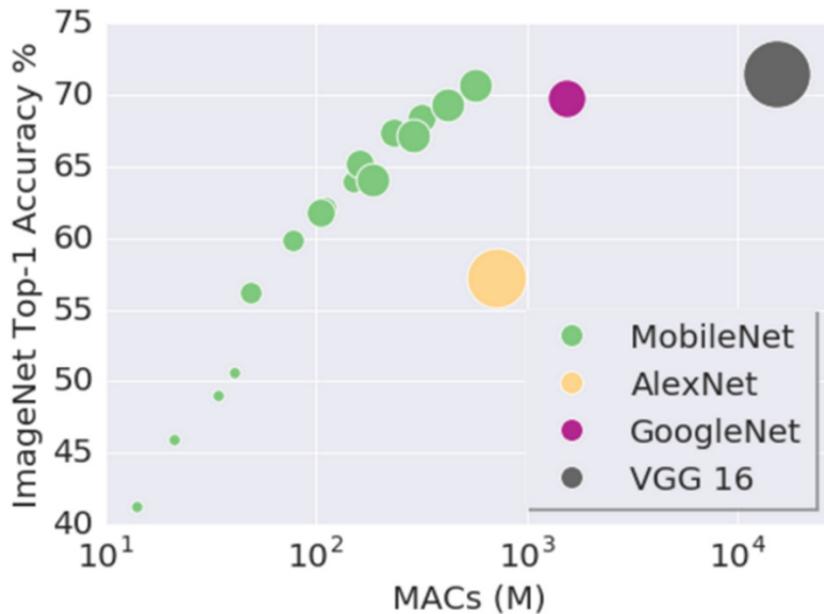


Table 9. Smaller MobileNet Comparison to Popular Models

| Model | ImageNet Accuracy | Million Mult-Adds | Million Parameters |
|--------------------|-------------------|-------------------|--------------------|
| 0.50 MobileNet-160 | 60.2% | 76 | 1.32 |
| SqueezeNet | 57.5% | 1700 | 1.25 |
| AlexNet | 57.2% | 720 | 60 |

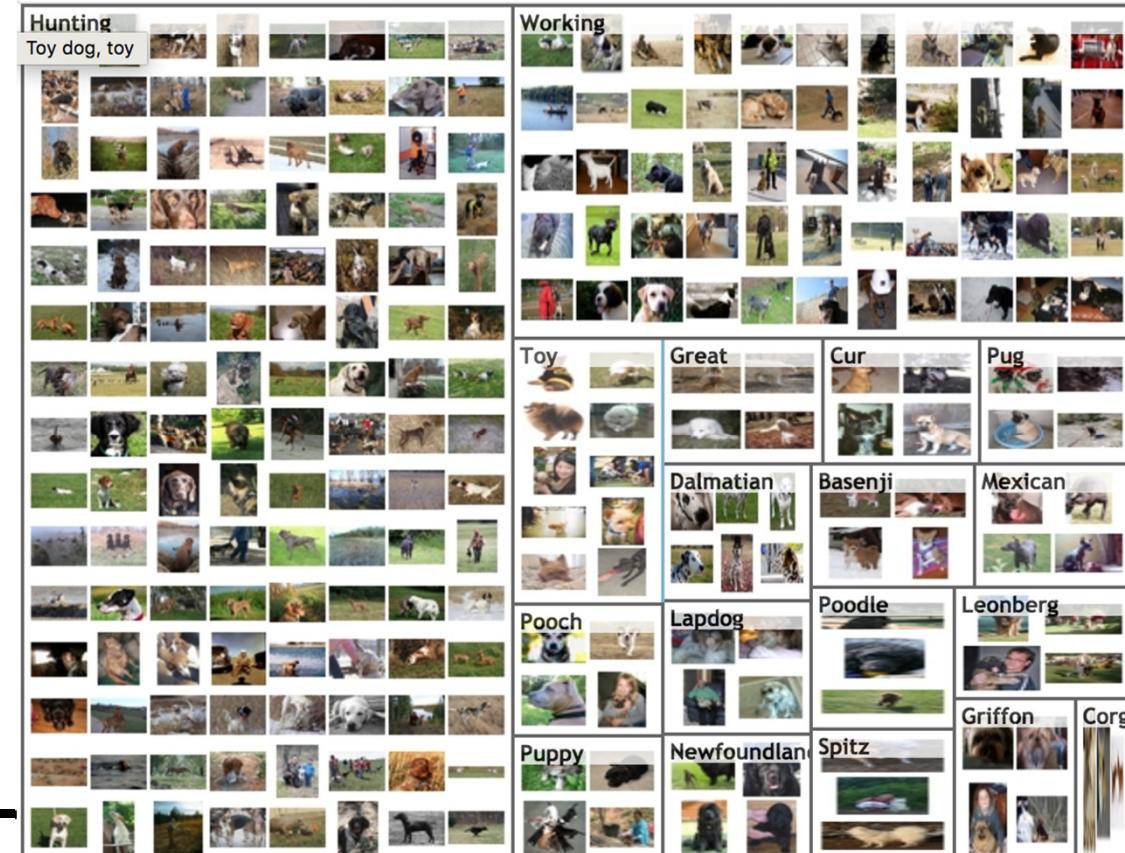
* Multiply-Accumulates (MACs) measures the number of fused Multiplication and Addition operations

ImageNet Dataset

[**ImageNet**](#): a **dataset of millions** of images with labels for **1000** different classes of **objects**, like dogs, cats, and fruits.

ImageNet 2011 Fall Release (32326)

- plant, flora, plant life (4486)
- geological formation, formation (175)
- natural object (1112)
- sport, athletics (176)
- artifact, artefact (10504)
- fungus (308)
- person, individual, someone, somebody, mortal, soul (6978)
- animal, animate being, beast, brute, creature, fauna (3998)
- Misc (20400)



ImageNet Dataset

List of 1000 different classes of objects

- 0: 'tench, *Tinca tinca*',
- 1: 'goldfish, *Carassius auratus*',
- 2: 'great white shark, white shark, man-eater, man-eating shark, ' + '*Carcharodon carcharias*',
- 3: 'tiger shark, *Galeocerdo cuvieri*',
- 4: 'hammerhead, hammerhead shark',
- 5: 'electric ray, crampfish, numbfish, torpedo',
- 6: 'stingray',
- 7: 'cock',
- 8: 'hen',
- 9: 'ostrich, *Struthio camelus*',
- 10: 'brambling, *Fringilla montifringilla*',

Transfer Learning

Use the **knowledge** gained while solving **one problem** and **applying** it to a **different** but related problem.

Start with a **pre-trained model** that are good at one task, lets you train far more **quickly** and with **less data** than if you were to train from scratch.

Transfer Learning

Use the knowledge gained while recognizing the **Imagenet classes** could apply when trying to recognize **user customized classes**.

Benefits:

- **Short** training time
- Needs very **little data**
- All in the **browser**.

Image from Webcam → [MobileNet](#)(Second to last layer) → [KNN Classifier](#) → Output
(High-level semantic features of the image)

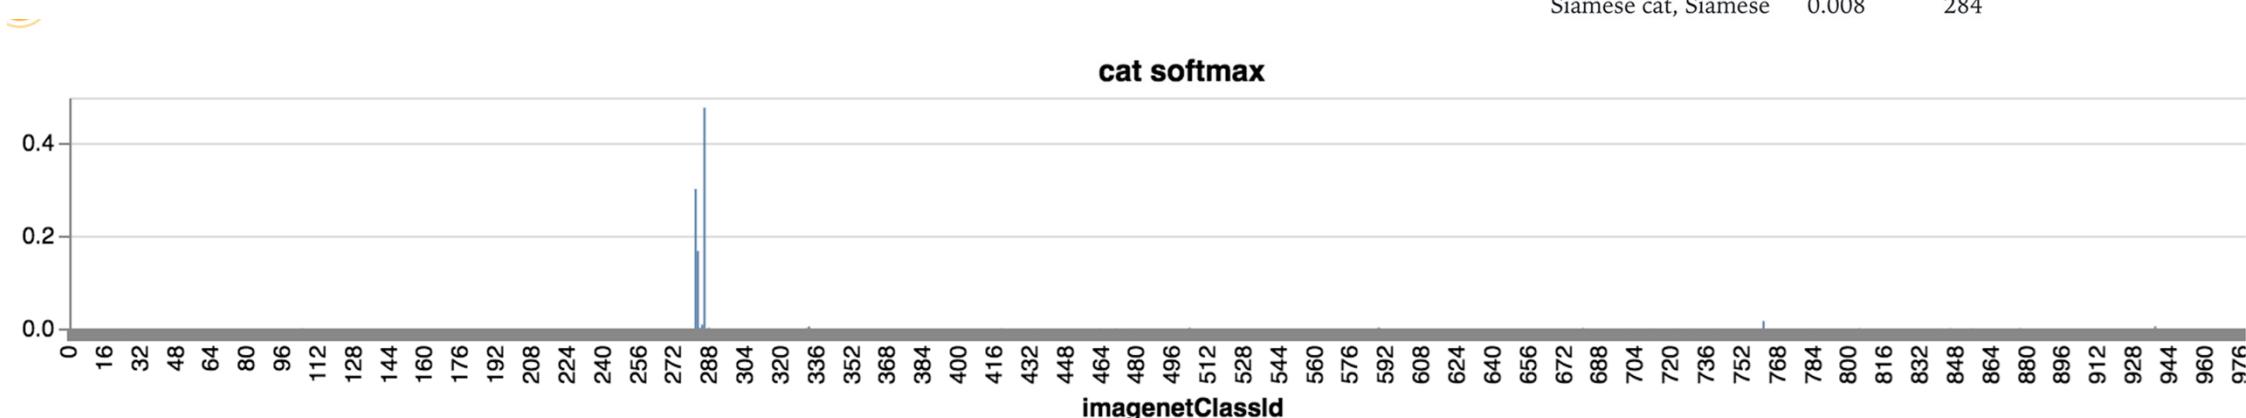
Output from MobileNet

Output from last layer of MobileNet:

- X axis: 0 - 999, 1000 ImageNet class name id
- Y axis: the probability (0 - 1) of the class



| class name | probability | imagenet class id |
|------------------------|-------------|-------------------|
| Egyptian cat | 0.478 | 285 |
| tabby, tabby cat | 0.300 | 281 |
| tiger cat | 0.167 | 282 |
| remote control, remote | 0.016 | 761 |
| Siamese cat, Siamese | 0.008 | 284 |



Source: [How to build a Teachable Machine with TensorFlow.js](https://www.tensorflow.org/tutorials/images/classification) by @nsthorat

Logits Activation from MobileNet

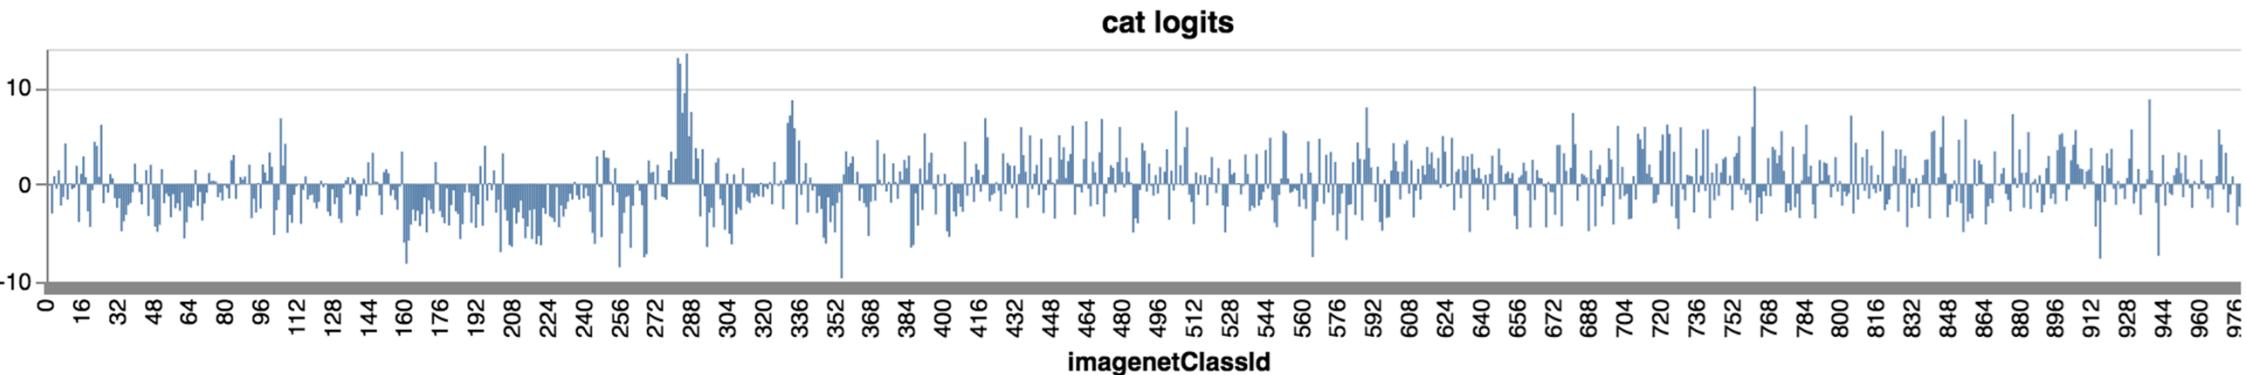
Activations: **output** of each layer

Softmax: calculates a **probability** for every possible class

Logits: **unnormalized** predictions vector



| class name | probability | imagenet class id |
|------------------------|-------------|-------------------|
| Egyptian cat | 0.478 | 285 |
| tabby, tabby cat | 0.300 | 281 |
| tiger cat | 0.167 | 282 |
| remote control, remote | 0.016 | 761 |
| Siamese cat, Siamese | 0.008 | 284 |

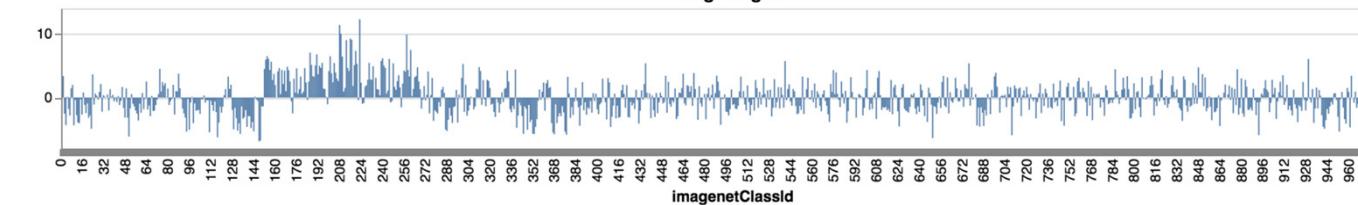
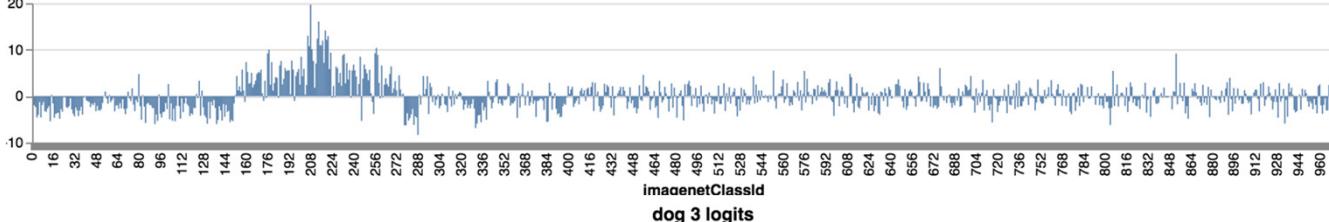
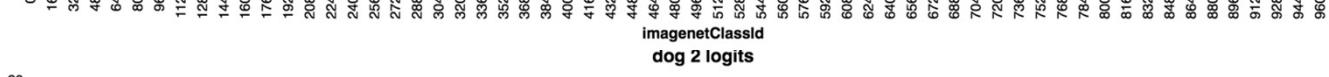
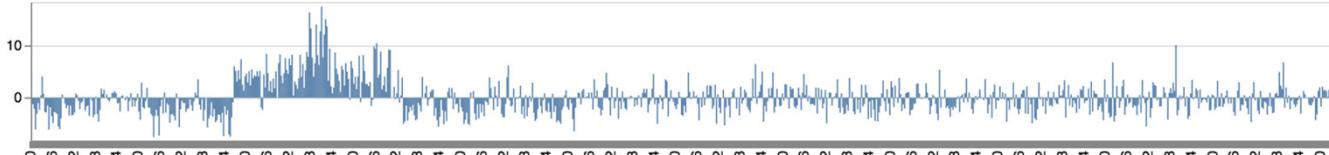


Source: [How to build a Teachable Machine with TensorFlow.js](#) by @nsthorat

More Logits Activation examples from MobileNet



dog 1 logits

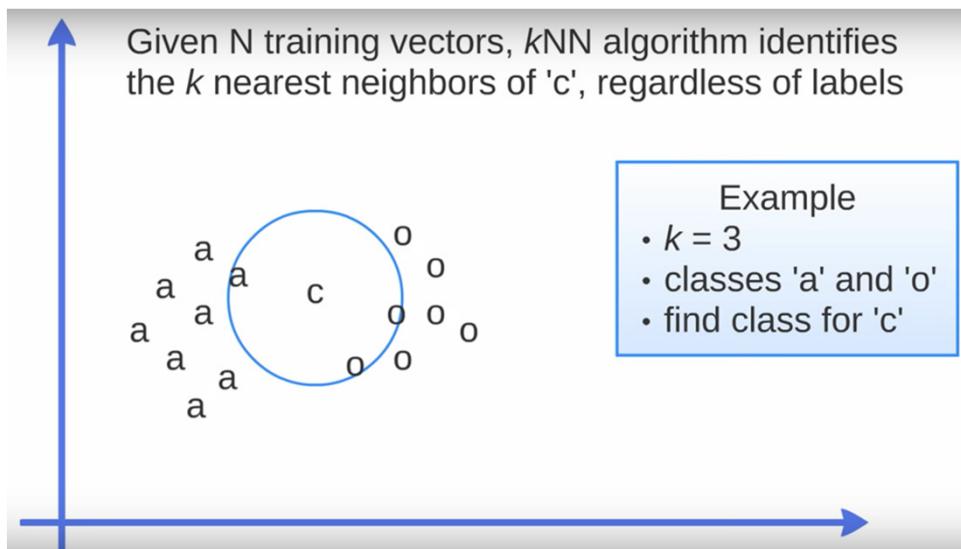


Source: [How to build a Teachable Machine with TensorFlow.js](#) by @nsthorat

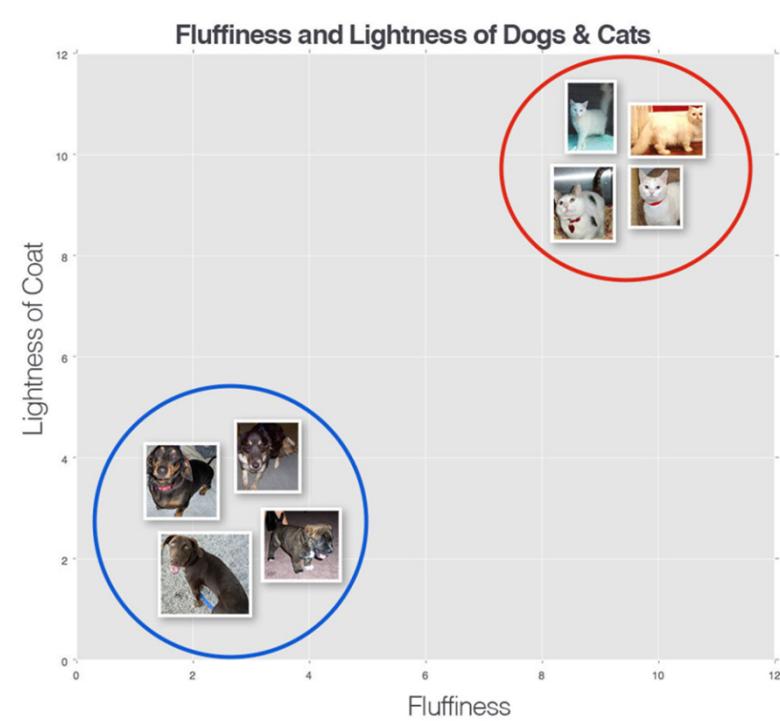
What is KNN Classifier?

K Nearest Neighbors, A classify algorithm, play with this interactive demo [here](#).

“Tell me who your neighbors are, and I’ll tell you who you are”



Source: [How kNN algorithm works](#), [K-NN classifier for image classification](#)



Build a KNN Image classifier with ml5

1. Data collection:

mobileNet.infer() → knnClassifier.addExample()

2. Prediction

mobileNet.infer() → knnClassifier.classify()

Image from Webcam → MobileNet(Second to last layer) → KNN Classifier → Output

Coding Session

- Teachable machine (image)
- KNN Image Classifier
- Regression with feature extractor

Code: <https://editor.p5js.org/ml5/sketches/KNNClassification> Video

Teachable Machine (/CNN Image Classifier)

Customizable image classifier on webcam images

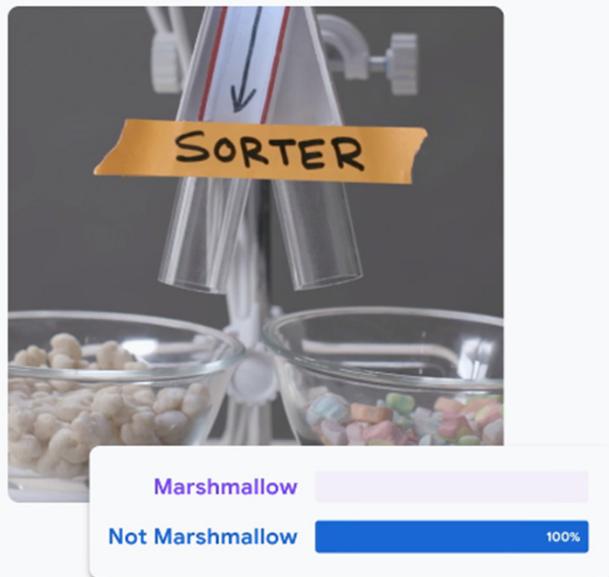
Teachable Machine

Train a computer to recognize your own images, sounds, & poses.

A fast, easy way to create machine learning models for your sites, apps, and more – no expertise or coding required.

[Get Started](#)





Projects related to Teachable Machine (Image Recognition)

- [Teachable Machine 2.0: Making AI easier for everyone](#)
- [Objectifier Spatial Programming](#)
- [Pacman](#)
- [Fruit quality detection](#)
- [Garbage detection](#)
- [Flower recognizer](#)
- [Teachable Snake](#)
- [Teachable Sorter](#)
- [Teachable Arcade](#)

Main Takeaways

1. What is MobileNet
2. What is KNN Classification?
3. How does Teachable Machine work?
4. What is Transfer Learning?