

# Object Detection

Notes based on  
CS231n, Stanford University, and  
EECS 498-007 / 598-005, University of Michigan  
with permission from [Justin Johnson](#)

# So far: Image Classification



This image is CC0 public domain




Figure copyright Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton, 2012. Reproduced with permission.

**Vector:**  
4096

## Fully-Connected: 4096 to 1000

## Class Scores

Cat: 0.9

Dog: 0.05

Car: 0.01

...


# Computer Vision Tasks

## Classification




**CAT**

## Semantic Segmentation



**GRASS, CAT, TREE, SKY**

## Object Detection



**DOG, DOG, CAT**

## Instance Segmentation



**DOG, DOG, CAT**

No spatial extent

No objects, just pixels

Multiple Objects

[This image is CC0 public domain](#)


# Today: Object Detection

## Classification




CAT

## Semantic Segmentation



GRASS, CAT, TREE,  
SKY

## Object Detection



DOG, DOG, CAT

## Instance Segmentation



DOG, DOG, CAT

No spatial extent

No objects, just pixels

Multiple Objects

[This image is CC0 public domain](#)




# Object Detection: Task Definition

**Input:** Single RGB Image


**Output:** A set of detected objects;  
For each object predict:

1. Category label (from fixed, known set of categories)
2. Bounding box (four numbers: x, y, width, height)



# Object Detection: Challenges


- **Multiple outputs:** Need to output variable numbers of objects per image
- **Multiple types of output:** Need to predict "what" (category label) as well as "where" (bounding box)
- **Large images:** Classification works at 224x224; need higher resolution for detection, often ~800x600



# Detecting a single object



[This image](#) is [CC0 public domain](#)




**Vector:**  
4096

# Detecting a single object



[This image](#) is [CC0 public domain](#)



**Vector:**  
4096

Fully  
Connected:  
4096 to 1000

“What”

**Class Scores**

Cat: 0.9  
Dog: 0.05  
Car: 0.01  
...

**Correct label:**

Cat

↓  
**Softmax  
Loss**




# Detecting a single object



[This image](#) is [CC0 public domain](#)

Treat localization as a regression problem!



**Vector:**  
4096

“Where”

Fully  
Connected:  
4096 to 1000

“What”

**Class Scores**

Cat: 0.9  
Dog: 0.05  
Car: 0.01  
...

**Correct label:**  
Cat

**Softmax  
Loss**

Fully  
Connected:  
4096 to 4

**Box  
Coordinates**  
(x, y, w, h)

**L2 Loss**


**Correct box:**  
(x', y', w', h')

# Detecting a single object



[This image](#) is [CC0 public domain](#)

Treat localization as a regression problem!



**Vector:**  
4096

“Where”

Fully  
Connected:  
4096 to 1000

“What”

**Class Scores**

Cat: 0.9  
Dog: 0.05  
Car: 0.01  
...

**Correct label:**

Cat

**Softmax  
Loss**

**Weighted  
Sum**

**Loss**

Fully  
Connected:  
4096 to 4

**Box  
Coordinates**  
(x, y, w, h)

**L2 Loss**


**Correct box:**  
(x', y', w', h')

# Detecting a single object



This image is [CC0 public domain](#)

Treat localization as a regression problem!



**Vector:**  
4096

“Where”

Fully  
Connected:  
4096 to 1000

“What”

**Class Scores**

Cat: 0.9  
Dog: 0.05  
Car: 0.01  
...

Fully  
Connected:  
4096 to 4

**Box  
Coordinates**  
(x, y, w, h)

**Correct label:**  
Cat

**Softmax  
Loss**

Multitask  
Loss

**Weighted  
Sum**


**Loss**

**L2 Loss**

**Correct box:**  
(x', y', w', h')

# Detecting a single object

Often pretrained  
on ImageNet  
(Transfer learning)



[This image](#) is [CC0 public domain](#)

Treat localization as a  
regression problem!

Fully  
Connected:  
4096 to 1000

“What”

**Class Scores**

Cat: 0.9  
Dog: 0.05  
Car: 0.01  
...

**Correct label:**

Cat

**Softmax  
Loss**

Multitask  
Loss

**Weighted  
Sum**

**Loss**

Fully  
Connected:  
4096 to 4

**Box  
Coordinates**  
(x, y, w, h)

**L2 Loss**

**Correct box:**  
(x', y', w', h')


“Where”

# Detecting a single object

Often pretrained  
on ImageNet  
(Transfer learning)



[This image is CC0 public domain](#)



**Vector:**  
4096

Treat localization as a  
regression problem!

“Where”

“What”

Fully  
Connected:  
4096 to 1000

**Class Scores**

Cat: 0.9  
Dog: 0.05  
Car: 0.01  
...

**Correct label:**  
Cat

**Softmax  
Loss**

Multitask  
Loss

**Weighted  
Sum**

**Loss**

Fully  
Connected:  
4096 to 4

**Box  
Coordinates**  
(x, y, w, h)

**L2 Loss**


**Correct box:**  
(x', y', w', h')

**Problem:** Images can have  
more than one object!




# Detecting Multiple Objects

Need different numbers  
of outputs per image



CAT:  $(x, y, w, h)$

4 numbers




DOG:  $(x, y, w, h)$

DOG:  $(x, y, w, h)$

CAT:  $(x, y, w, h)$

12 numbers



DUCK:  $(x, y, w, h)$


DUCK:  $(x, y, w, h)$

....

Many  
numbers!

# Detecting Multiple Objects: Sliding Window

Apply a CNN to many different crops of the image, CNN classifies each crop as object or background





Dog? **NO**

Cat? **NO**

Background? **YES**

# Detecting Multiple Objects: Sliding Window

Apply a CNN to many different crops of the image, CNN classifies each crop as object or background




Dog? YES

Cat? NO

Background? NO

# Detecting Multiple Objects: Sliding Window

Apply a CNN to many different crops of the image, CNN classifies each crop as object or background



Dog? **YES**


Cat? **NO**

Background? **NO**



# Detecting Multiple Objects: Sliding Window

Apply a CNN to many different crops of the image, CNN classifies each crop as object or background



Dog? **NO**

Cat? **YES**

Background? **NO**



# Detecting Multiple Objects: Sliding Window

Apply a CNN to many different crops of the image, CNN classifies each crop as object or background

**Question:** How many possible boxes are there in an image of size  $H \times W$ ?



# Detecting Multiple Objects: Sliding Window

Apply a CNN to many different crops of the image, CNN classifies each crop as object or background



**Question:** How many possible boxes are there in an image of size  $H \times W$ ?

Consider a box of size  $h \times w$ :

Possible x positions:  $W - w + 1$

Possible y positions:  $H - h + 1$

Possible positions:

$$(W - w + 1) * (H - h + 1)$$

# Detecting Multiple Objects: Sliding Window

Apply a CNN to many different crops of the image, CNN classifies each crop as object or background



**Question:** How many possible boxes are there in an image of size  $H \times W$ ?

Consider a box of size  $h \times w$ :

Possible x positions:  $W - w + 1$

Possible y positions:  $H - h + 1$

Possible positions:

$(W - w + 1) * (H - h + 1)$

Total possible boxes:

$$\sum_{h=1}^H \sum_{w=1}^W (W - w + 1)(H - h + 1)$$

$$= \frac{H(H + 1)}{2} \frac{W(W + 1)}{2}$$

# Detecting Multiple Objects: Sliding Window

Apply a CNN to many different crops of the image, CNN classifies each crop as object or background

800 x 600 image  
has ~58M boxes!  
No way we can  
evaluate them all



**Question:** How many possible boxes are there in an image of size  $H \times W$ ?

Consider a box of size  $h \times w$ :

Possible x positions:  $W - w + 1$

Possible y positions:  $H - h + 1$

Possible positions:

$$(W - w + 1) * (H - h + 1)$$


Total possible boxes:

$$\sum_{h=1}^H \sum_{w=1}^W (W - w + 1)(H - h + 1)$$

$$= \frac{H(H + 1)}{2} \frac{W(W + 1)}{2}$$

# Region Proposals

- Find a small set of boxes that are likely to cover all objects
- Often based on heuristics: e.g. look for “blob-like” image regions
- Relatively fast to run; e.g. Selective Search gives 2000 region proposals in a few seconds on CPU



Alexe et al, “Measuring the objectness of image windows”, TPAMI 2012  
Uijlings et al, “Selective Search for Object Recognition”, IJCV 2013  
Cheng et al, “BING: Binarized normed gradients for objectness estimation at 300fps”, CVPR 2014  
Zitnick and Dollar, “Edge boxes: Locating object proposals from edges”, ECCV 2014



# R-CNN: Region-Based CNN

Input  
image




Girshick et al, "Rich feature hierarchies for accurate object detection and semantic segmentation", CVPR 2014.  
Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.

# R-CNN: Region-Based CNN



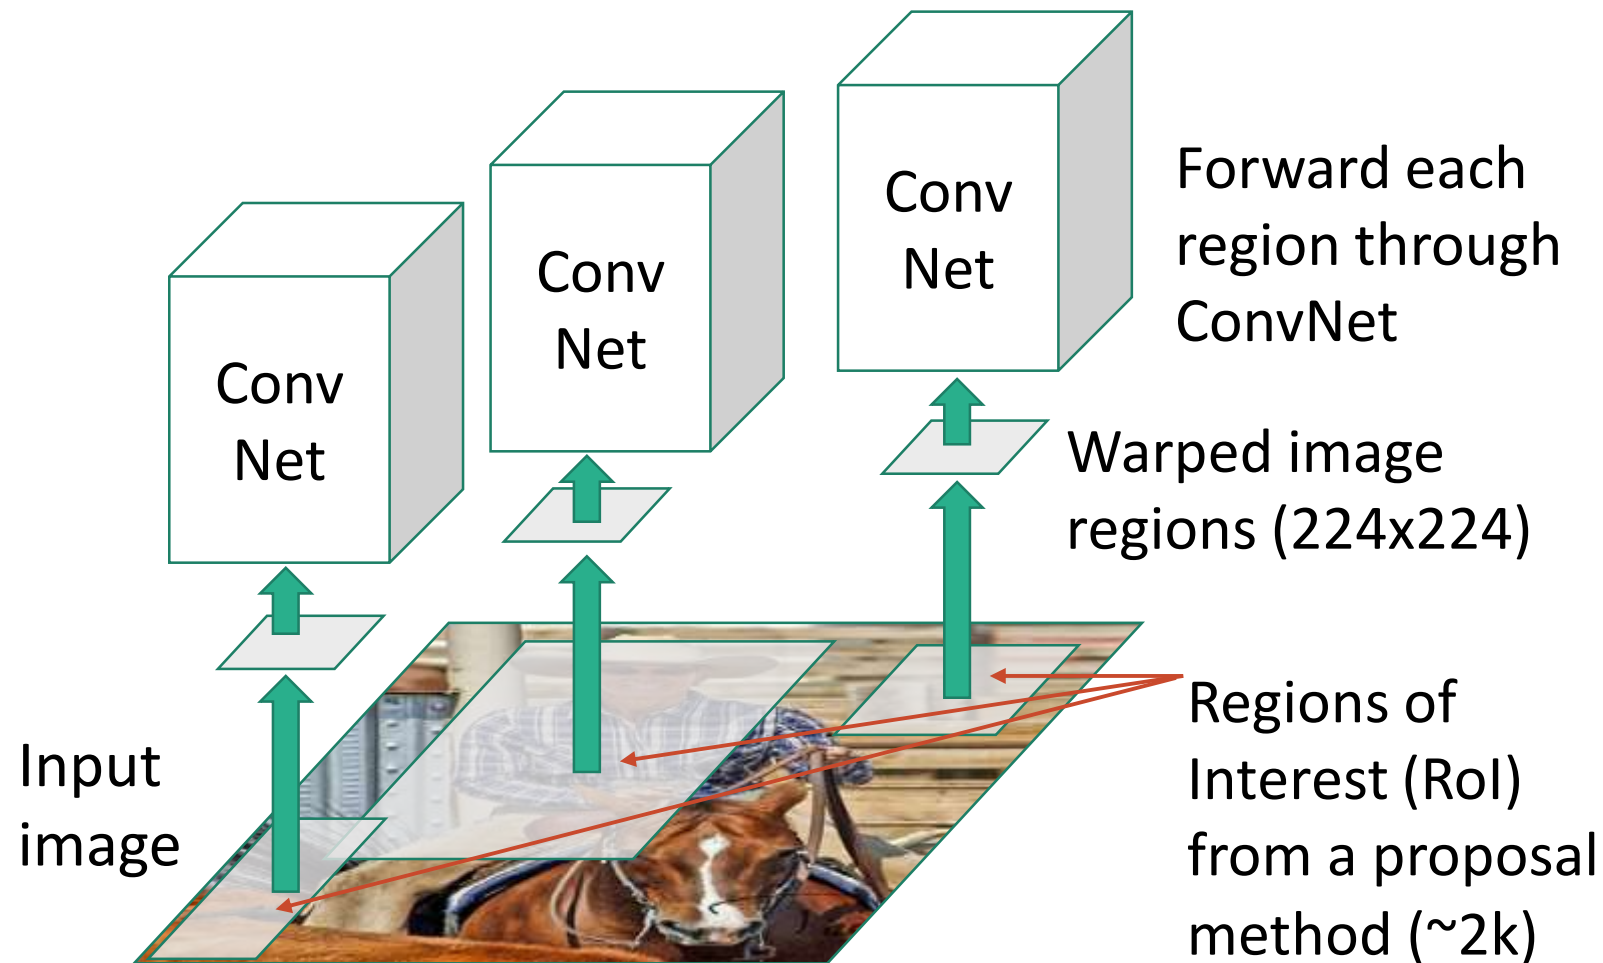
Girshick et al, "Rich feature hierarchies for accurate object detection and semantic segmentation", CVPR 2014.  
Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.

# R-CNN: Region-Based CNN



Girshick et al, "Rich feature hierarchies for accurate object detection and semantic segmentation", CVPR 2014.  
Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.

# R-CNN: Region-Based CNN

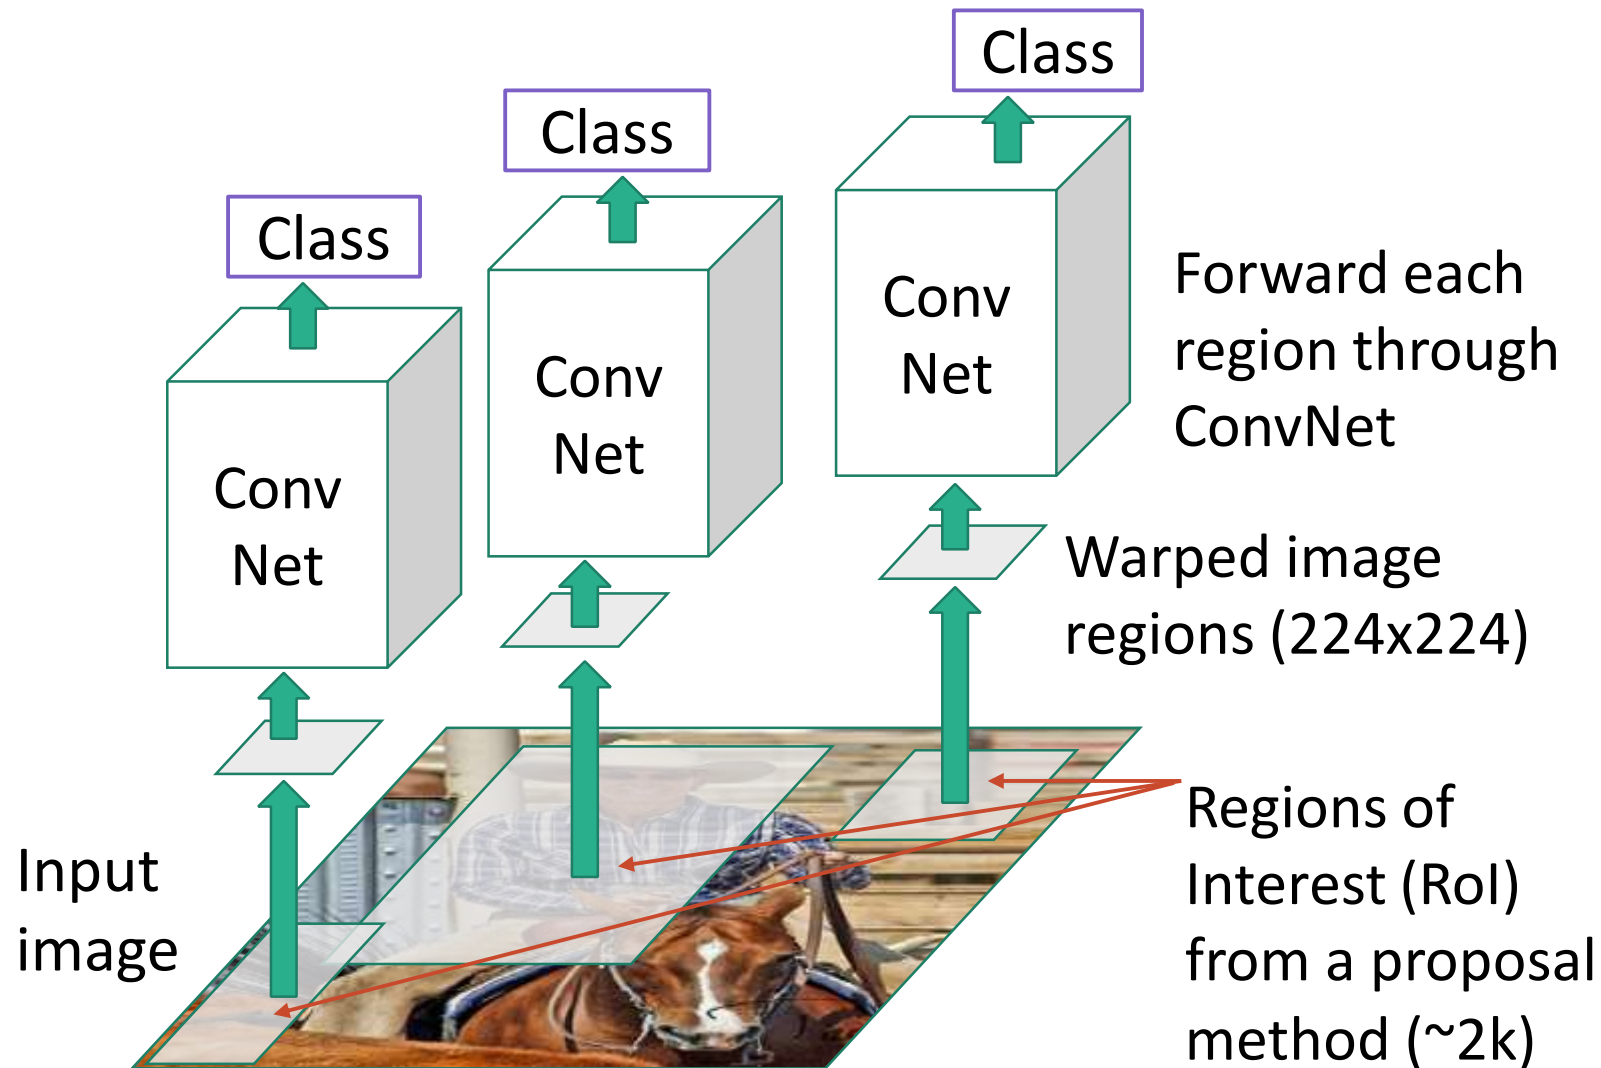


Girshick et al, "Rich feature hierarchies for accurate object detection and semantic segmentation", CVPR 2014.

Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.

# R-CNN: Region-Based CNN


Classify each region



Girshick et al, "Rich feature hierarchies for accurate object detection and semantic segmentation", CVPR 2014.  
Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.



# R-CNN: Region-Based CNN



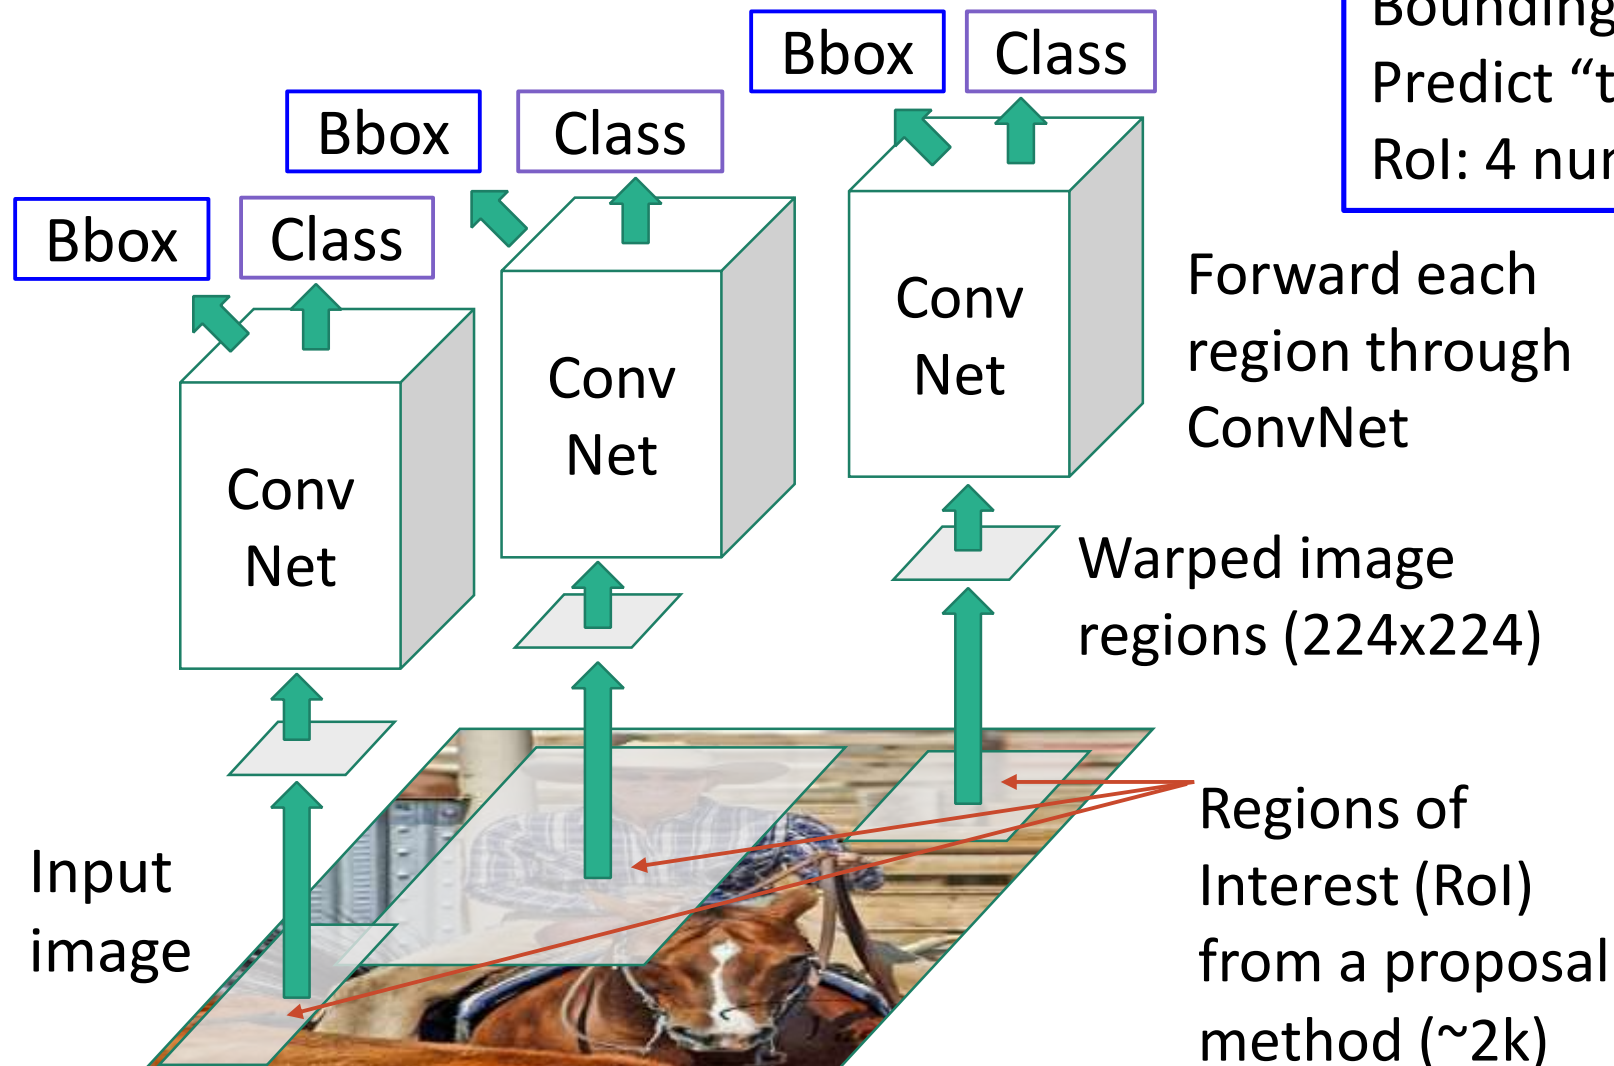
Classify each region

Bounding box regression:  
Predict "transform" to correct the  
RoI: 4 numbers ( $t_x, t_y, t_h, t_w$ )

Girshick et al, "Rich feature hierarchies for accurate object detection and semantic segmentation", CVPR 2014.

Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.

# R-CNN: Region-Based CNN



Classify each region

Bounding box regression:  
Predict “transform” to correct the  
RoI: 4 numbers ( $t_x, t_y, t_h, t_w$ )

Forward each  
region through  
ConvNet

Region proposal: ( $p_x, p_y, p_h, p_w$ )  
Transform: ( $t_x, t_y, t_h, t_w$ )  
Output box: ( $b_x, b_y, b_h, b_w$ )


Translate relative to box size:  
 $b_x = p_x + p_w t_x$        $b_y = p_y + p_h t_y$

Log-space scale transform:  
 $b_w = p_w \exp(t_w)$        $b_h = p_h \exp(t_h)$

Girshick et al, “Rich feature hierarchies for accurate object detection and semantic segmentation”, CVPR 2014.

Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.

# R-CNN: Region-Based CNN




Classify each region

Bounding box regression:  
Predict "transform" to correct the RoI: 4 numbers ( $t_x, t_y, t_h, t_w$ )

Girshick et al, "Rich feature hierarchies for accurate object detection and semantic segmentation", CVPR 2014.

Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.

# R-CNN: Region-Based CNN



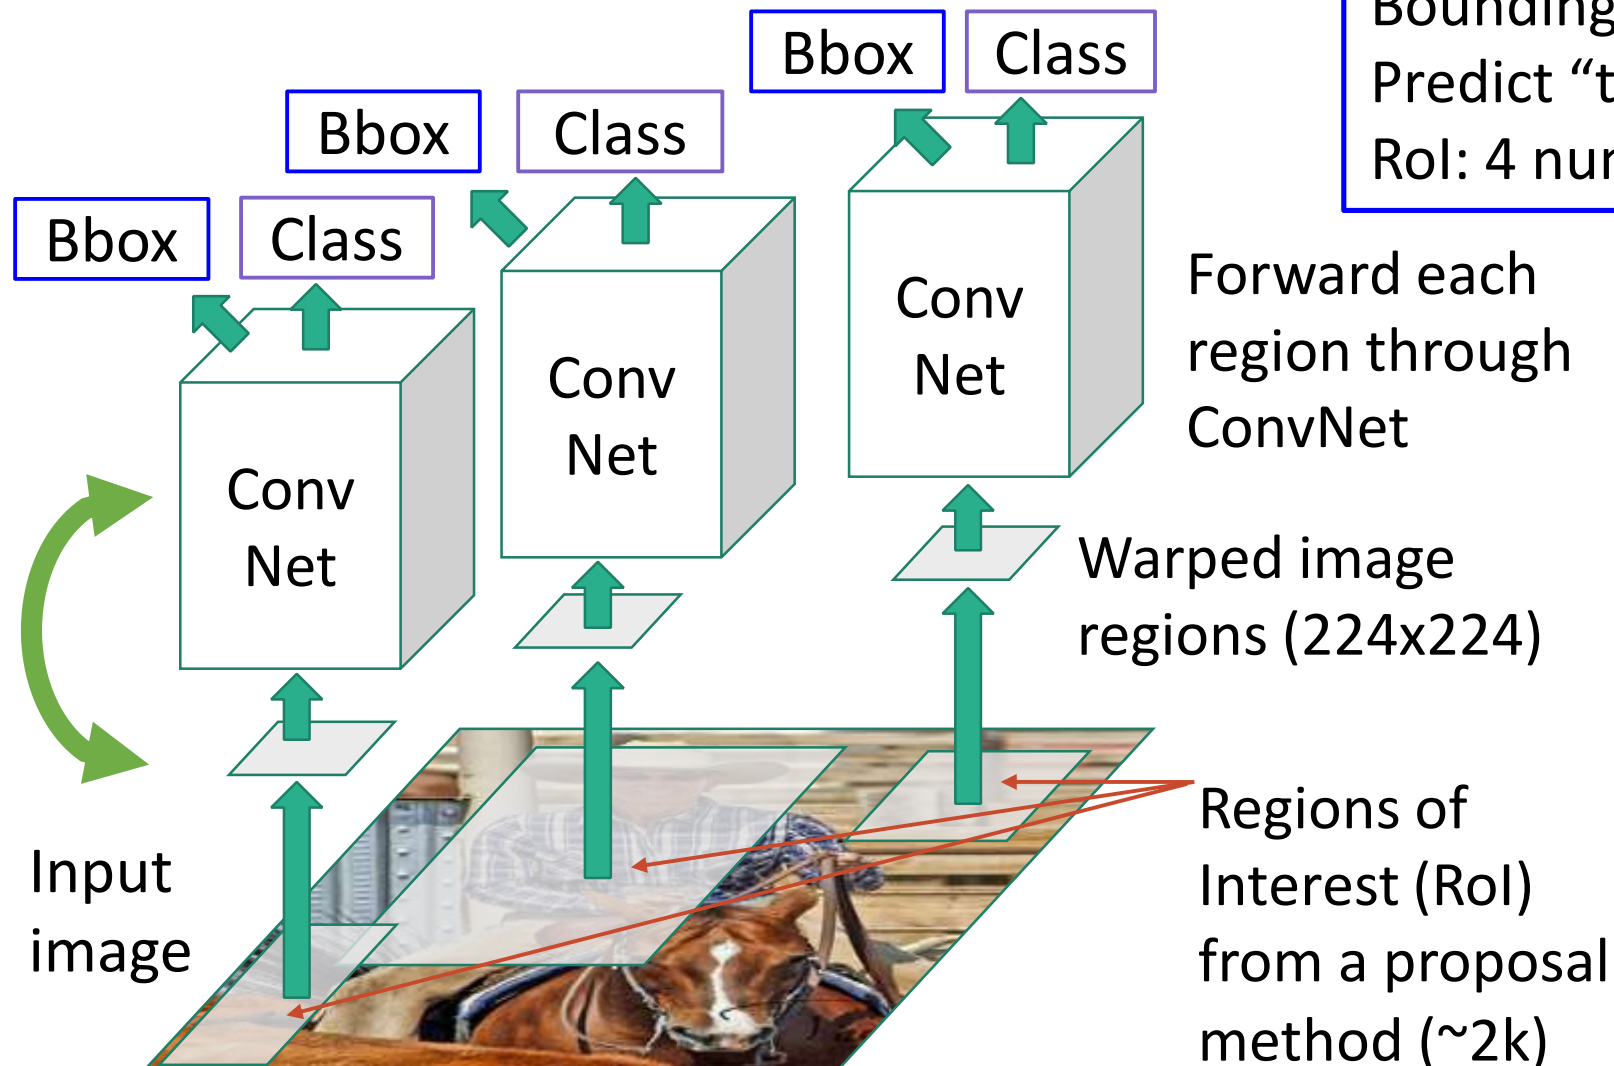
Classify each region

Bounding box regression:  
Predict "transform" to correct the  
RoI: 4 numbers ( $t_x, t_y, t_h, t_w$ )

**Problem: Very slow!**  
Need to do ~2k forward  
passes for each image!

Girshick et al, "Rich feature hierarchies for accurate object detection and semantic segmentation", CVPR 2014.  
Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.

# R-CNN: Region-Based CNN



Classify each region

Bounding box regression:  
Predict "transform" to correct the RoI: 4 numbers ( $t_x, t_y, t_h, t_w$ )

Forward each region through ConvNet

Warped image regions (224x224)

**Problem:** Very slow!  
Need to do ~2k forward passes for each image!

**Solution:** Run CNN  
\*before\* warping!

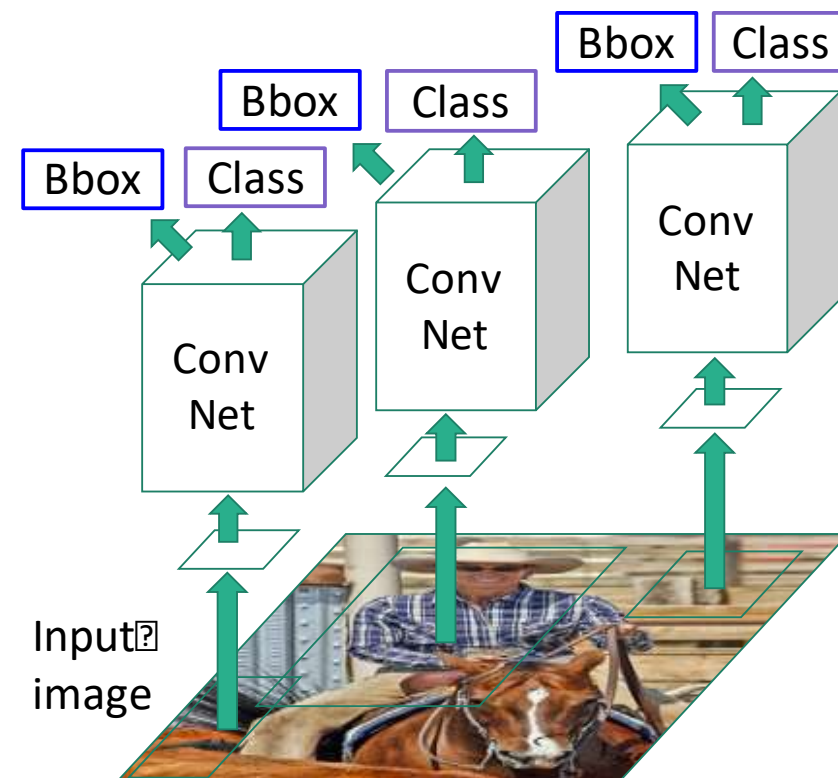
Girshick et al, "Rich feature hierarchies for accurate object detection and semantic segmentation", CVPR 2014.

Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.



## “Slow” R-CNN

Process each region  
independently




# Fast R-CNN



Input image

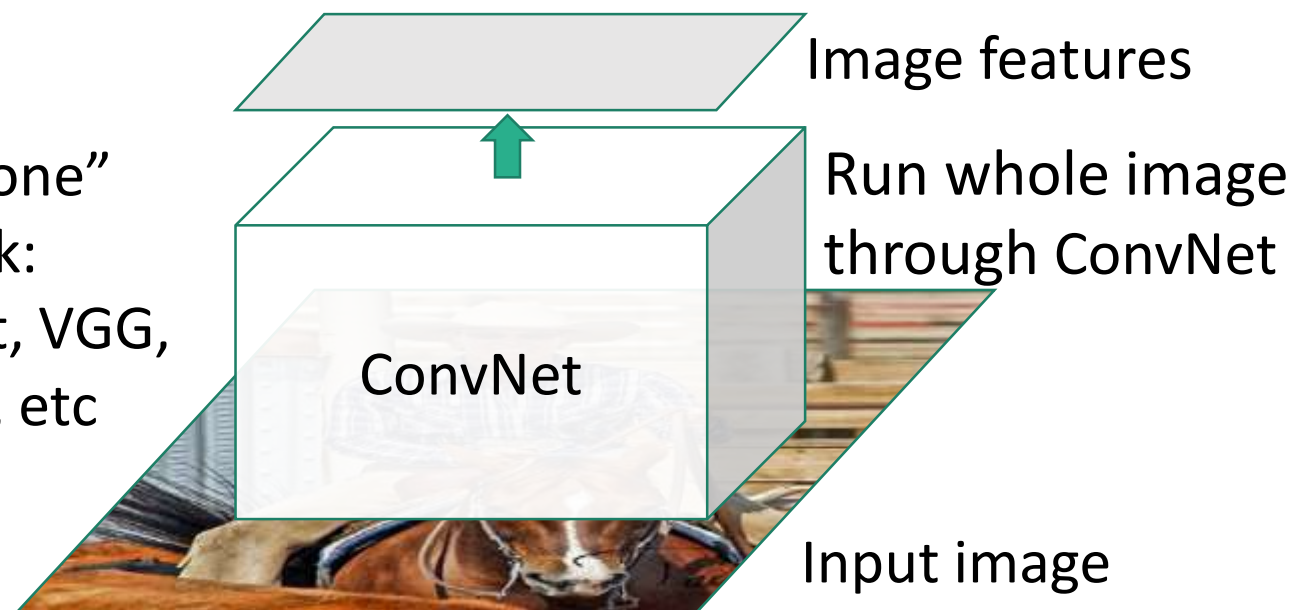
“Slow” R-CNN  
Process each region  
independently




Girshick, “Fast R-CNN”, ICCV 2015. Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.

# Fast R-CNN

“Backbone”  
network:  
AlexNet, VGG,  
ResNet, etc



“Slow” R-CNN  
Process each region  
independently

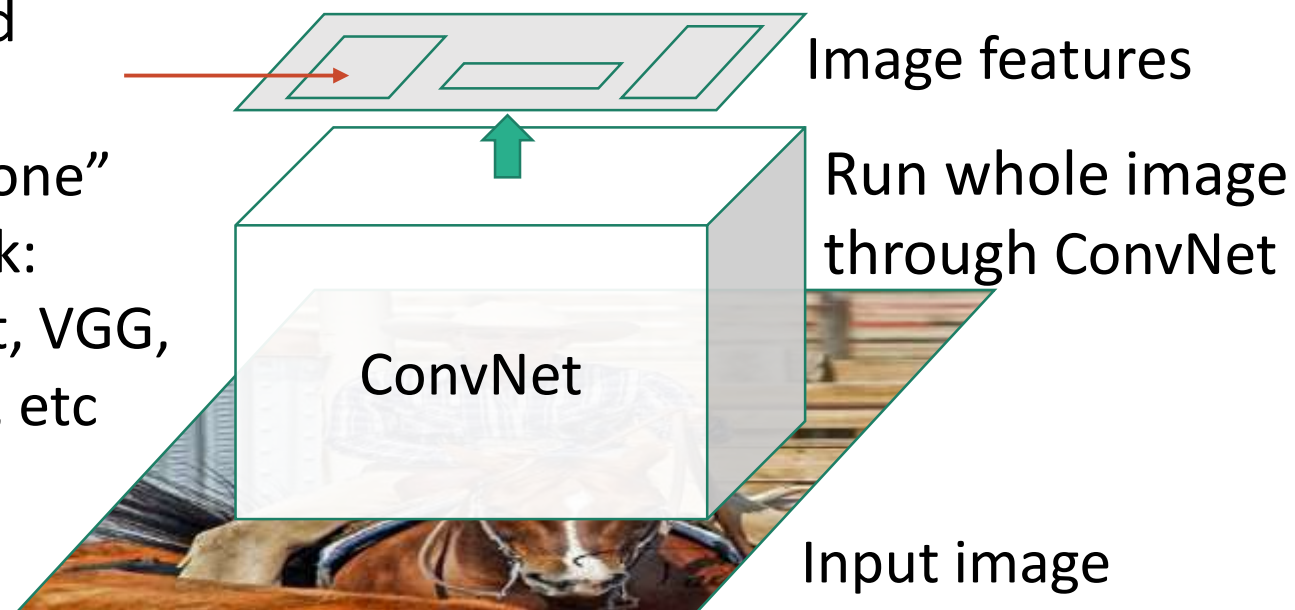


Girshick, “Fast R-CNN”, ICCV 2015. Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.

# Fast R-CNN


Regions of Interest (ROIs) from a proposal method

“Backbone” network:  
AlexNet, VGG, ResNet, etc



## “Slow” R-CNN

Process each region independently

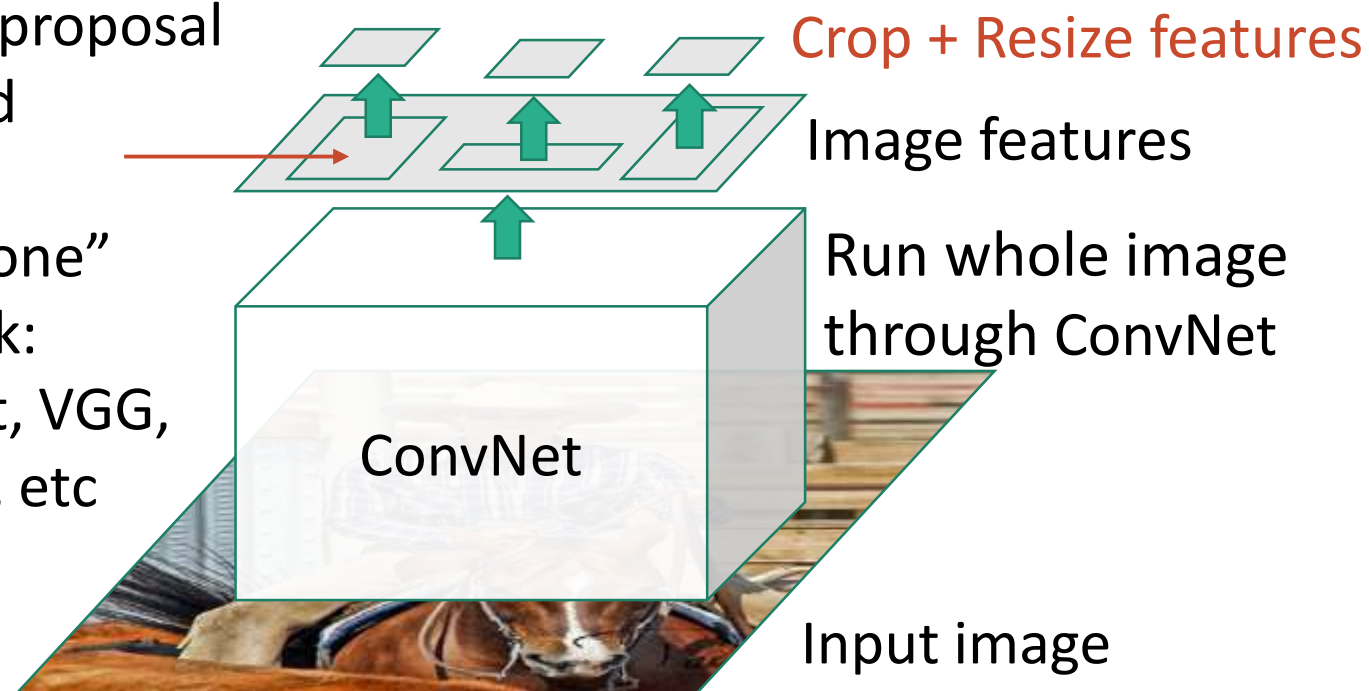


Girshick, “Fast R-CNN”, ICCV 2015. Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.


# Fast R-CNN

Regions of Interest (ROIs) from a proposal method

“Backbone” network:  
AlexNet, VGG, ResNet, etc



“Slow” R-CNN  
Process each region independently



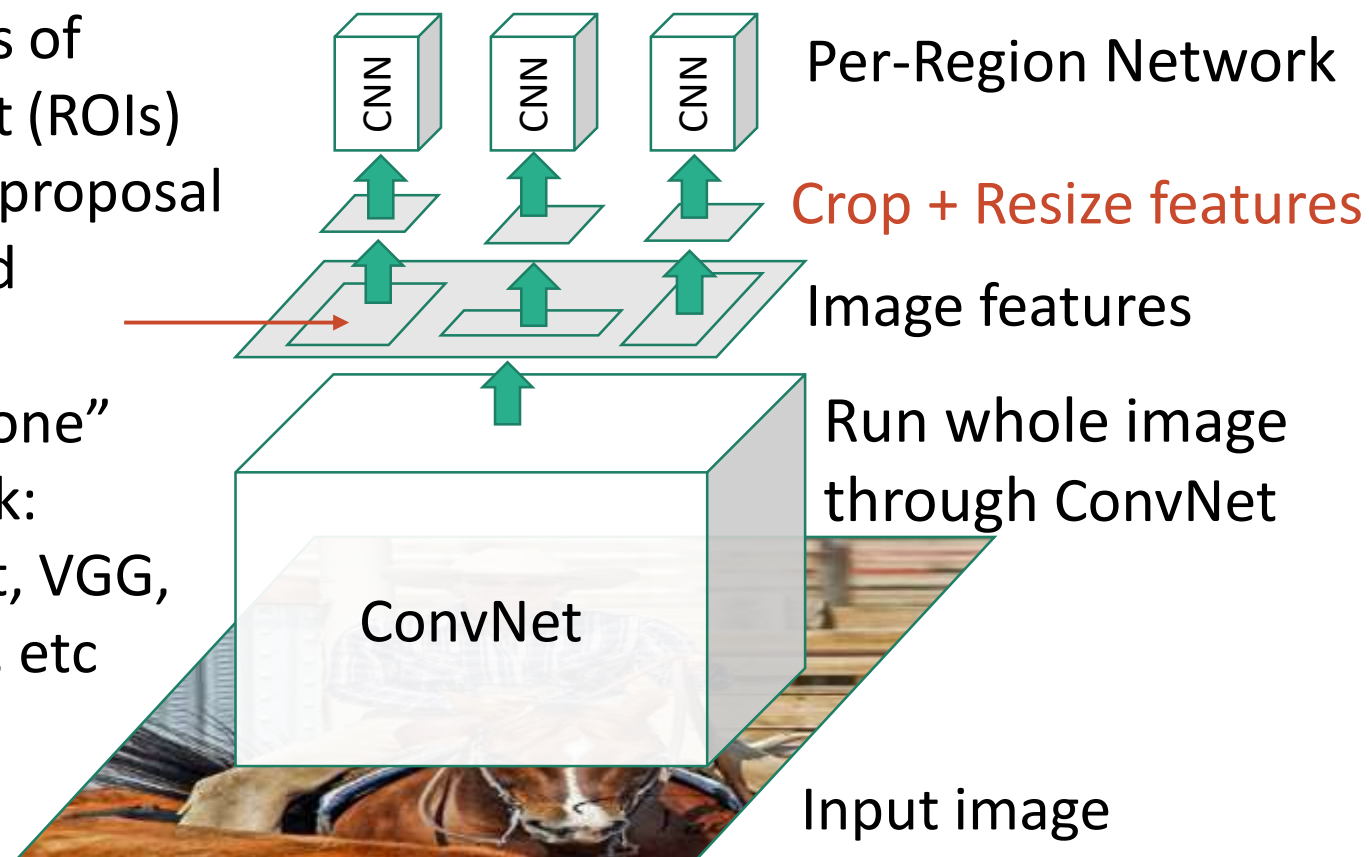
Girshick, “Fast R-CNN”, ICCV 2015. Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.



# Fast R-CNN

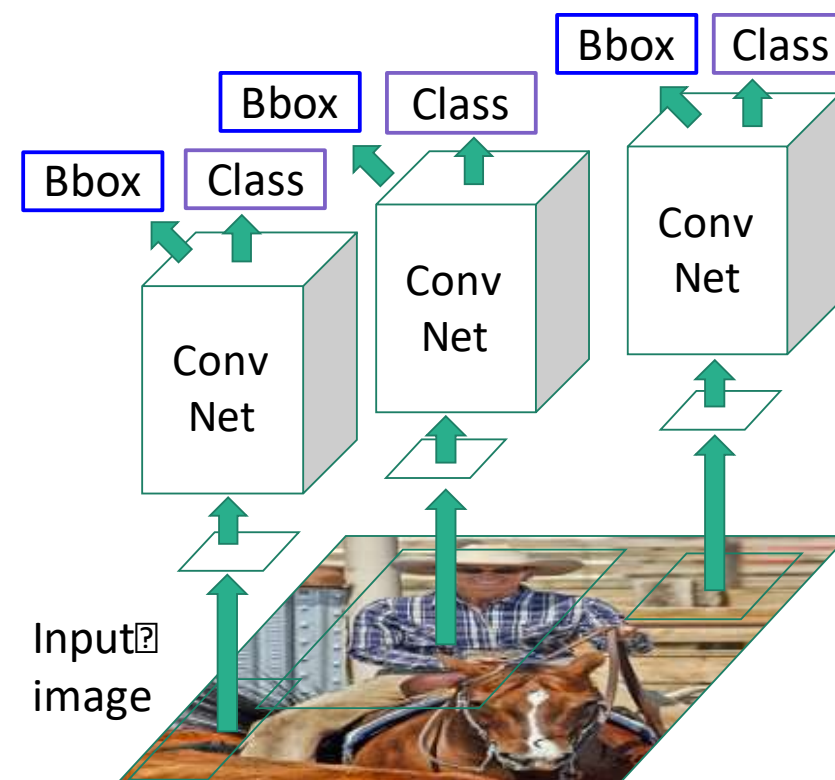
Regions of Interest (ROIs) from a proposal method

“Backbone” network:  
AlexNet, VGG, ResNet, etc

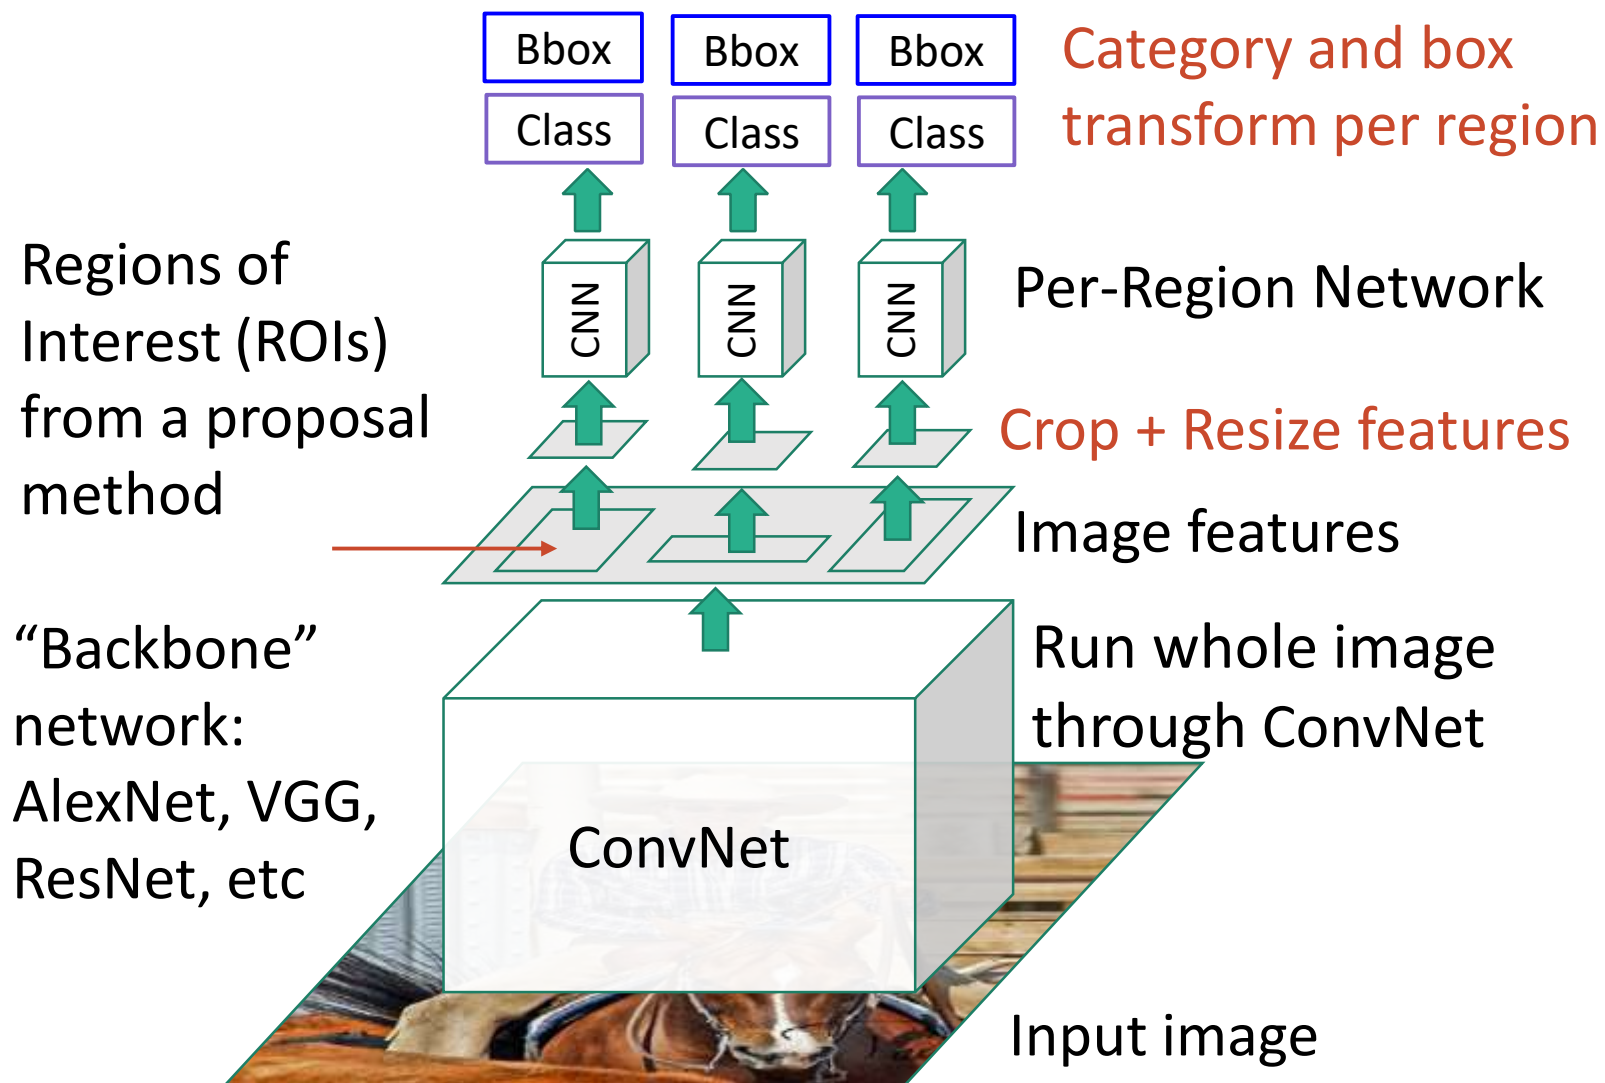


## “Slow” R-CNN

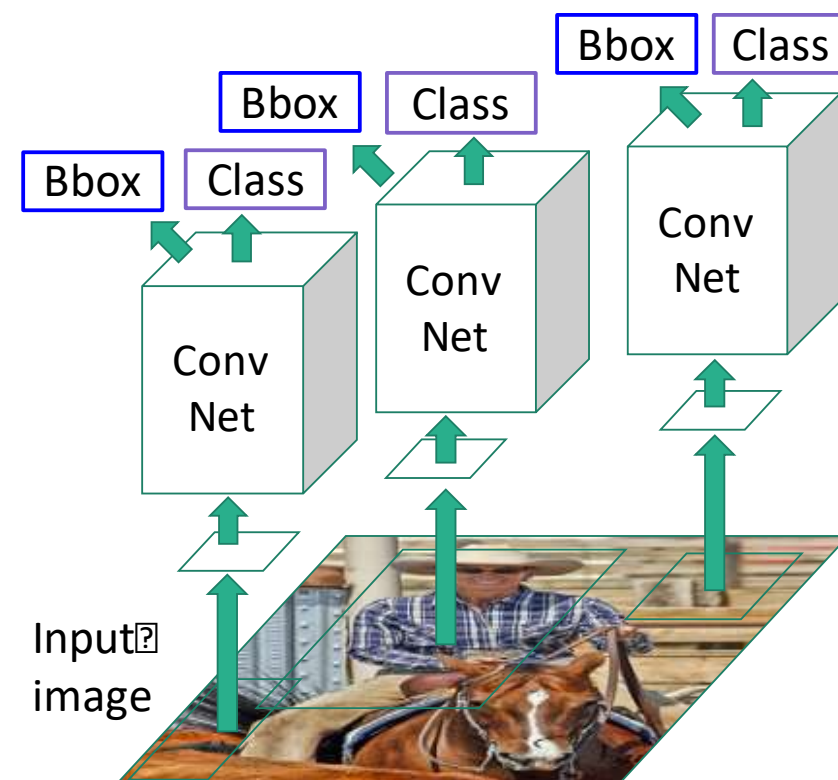
Process each region independently



# Fast R-CNN

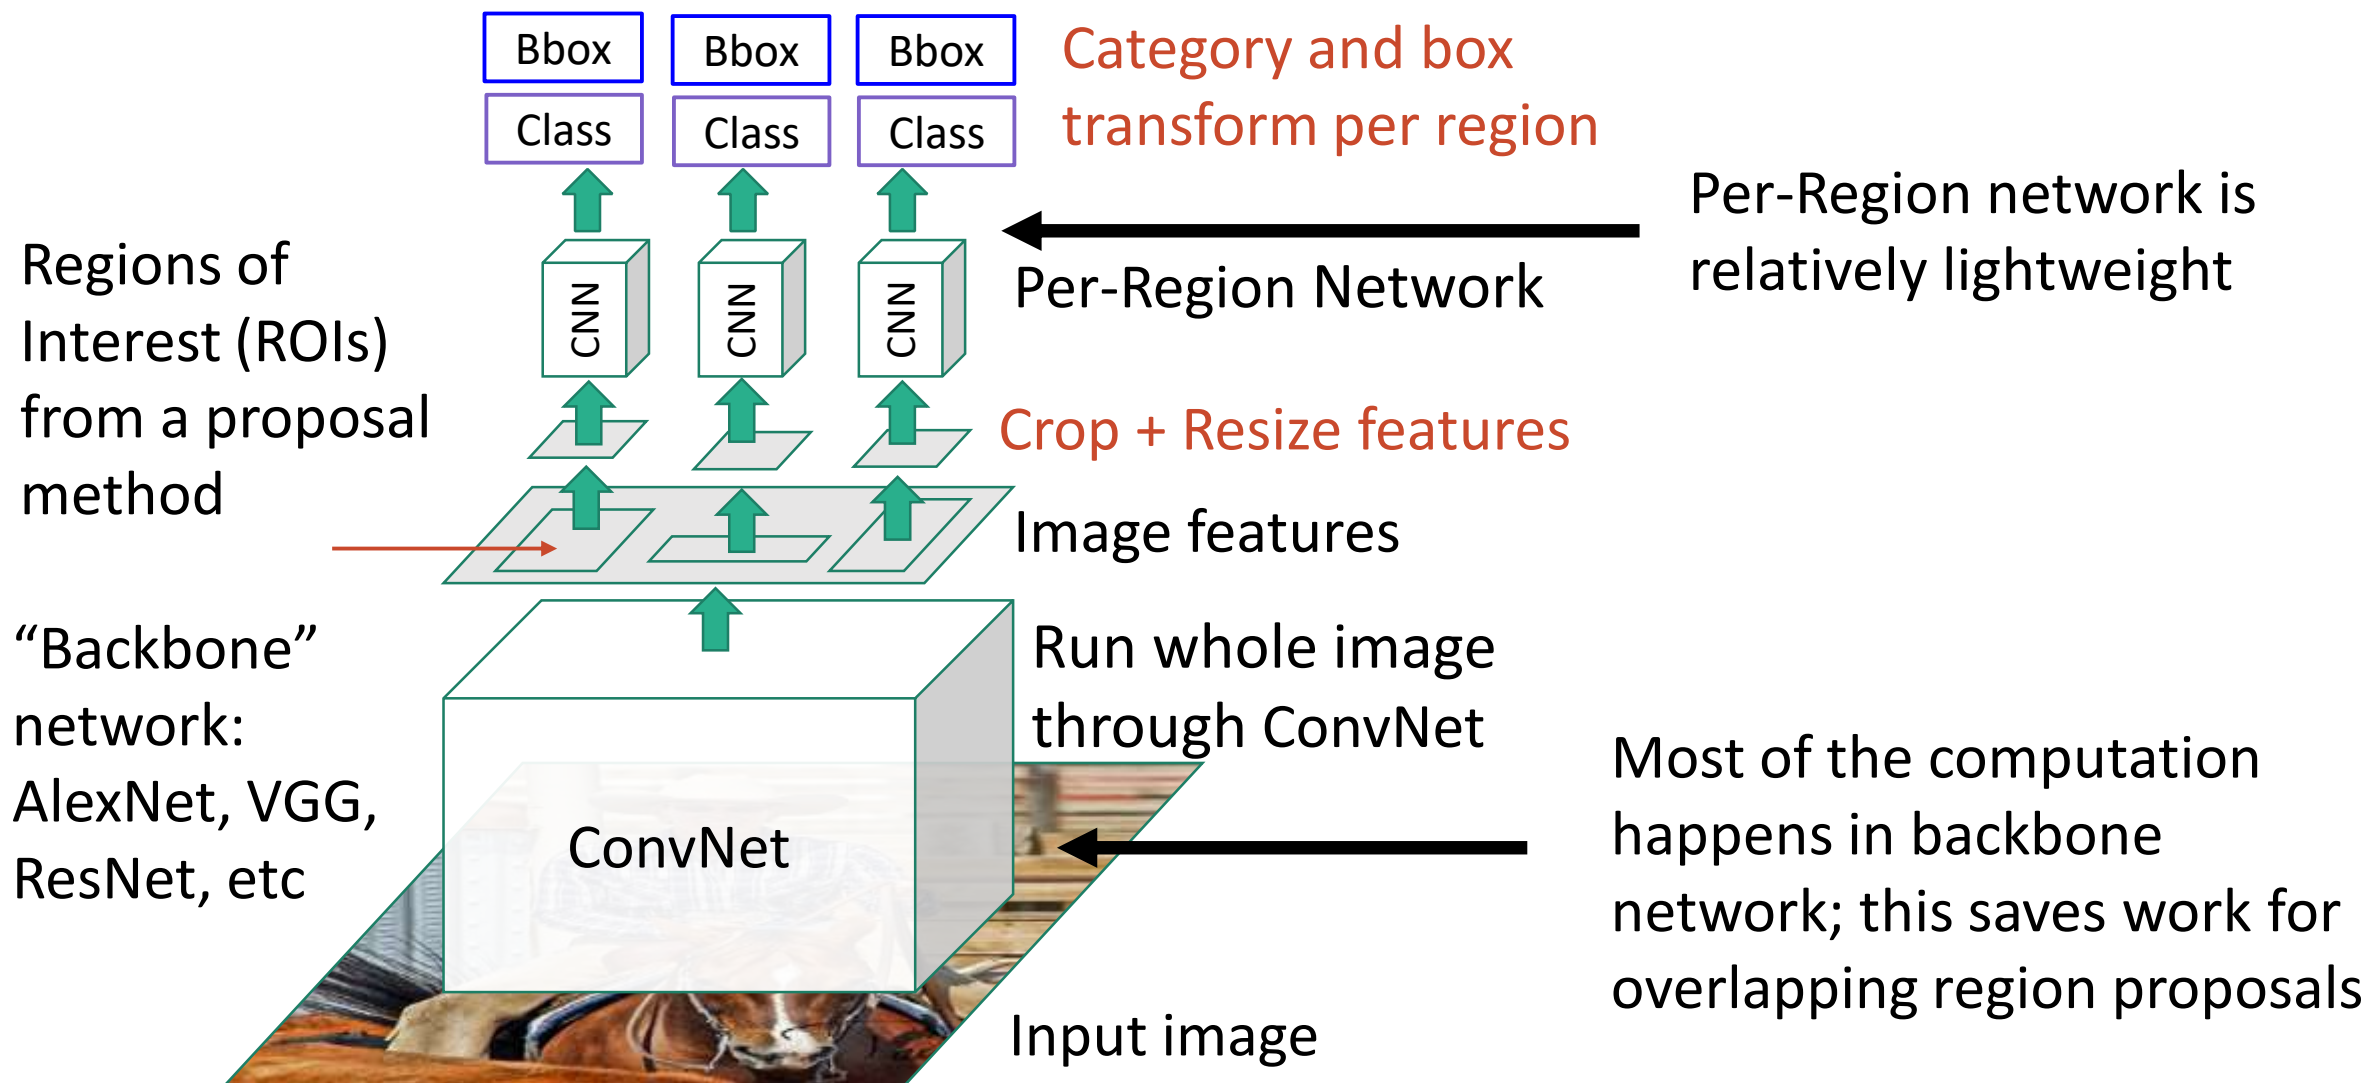


“Slow” R-CNN  
Process each region independently



Girshick, “Fast R-CNN”, ICCV 2015. Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.


# Fast R-CNN




Girshick, “Fast R-CNN”, ICCV 2015. Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.

# Fast R-CNN vs “Slow” R-CNN

## Training time (Hours)



## Test time (seconds)




Girshick et al, “Rich feature hierarchies for accurate object detection and semantic segmentation”, CVPR 2014.

He et al, “Spatial pyramid pooling in deep convolutional networks for visual recognition”, ECCV 2014


Girshick, “Fast R-CNN”, ICCV 2015

# Fast R-CNN vs “Slow” R-CNN

## Training time (Hours)



## Test time (seconds)



**Problem:** Runtime dominated by region proposals!

Girshick et al, “Rich feature hierarchies for accurate object detection and semantic segmentation”, CVPR 2014.


He et al, “Spatial pyramid pooling in deep convolutional networks for visual recognition”, ECCV 2014

Girshick, “Fast R-CNN”, ICCV 2015




# Fast R-CNN vs “Slow” R-CNN

## Training time (Hours)



## Test time (seconds)



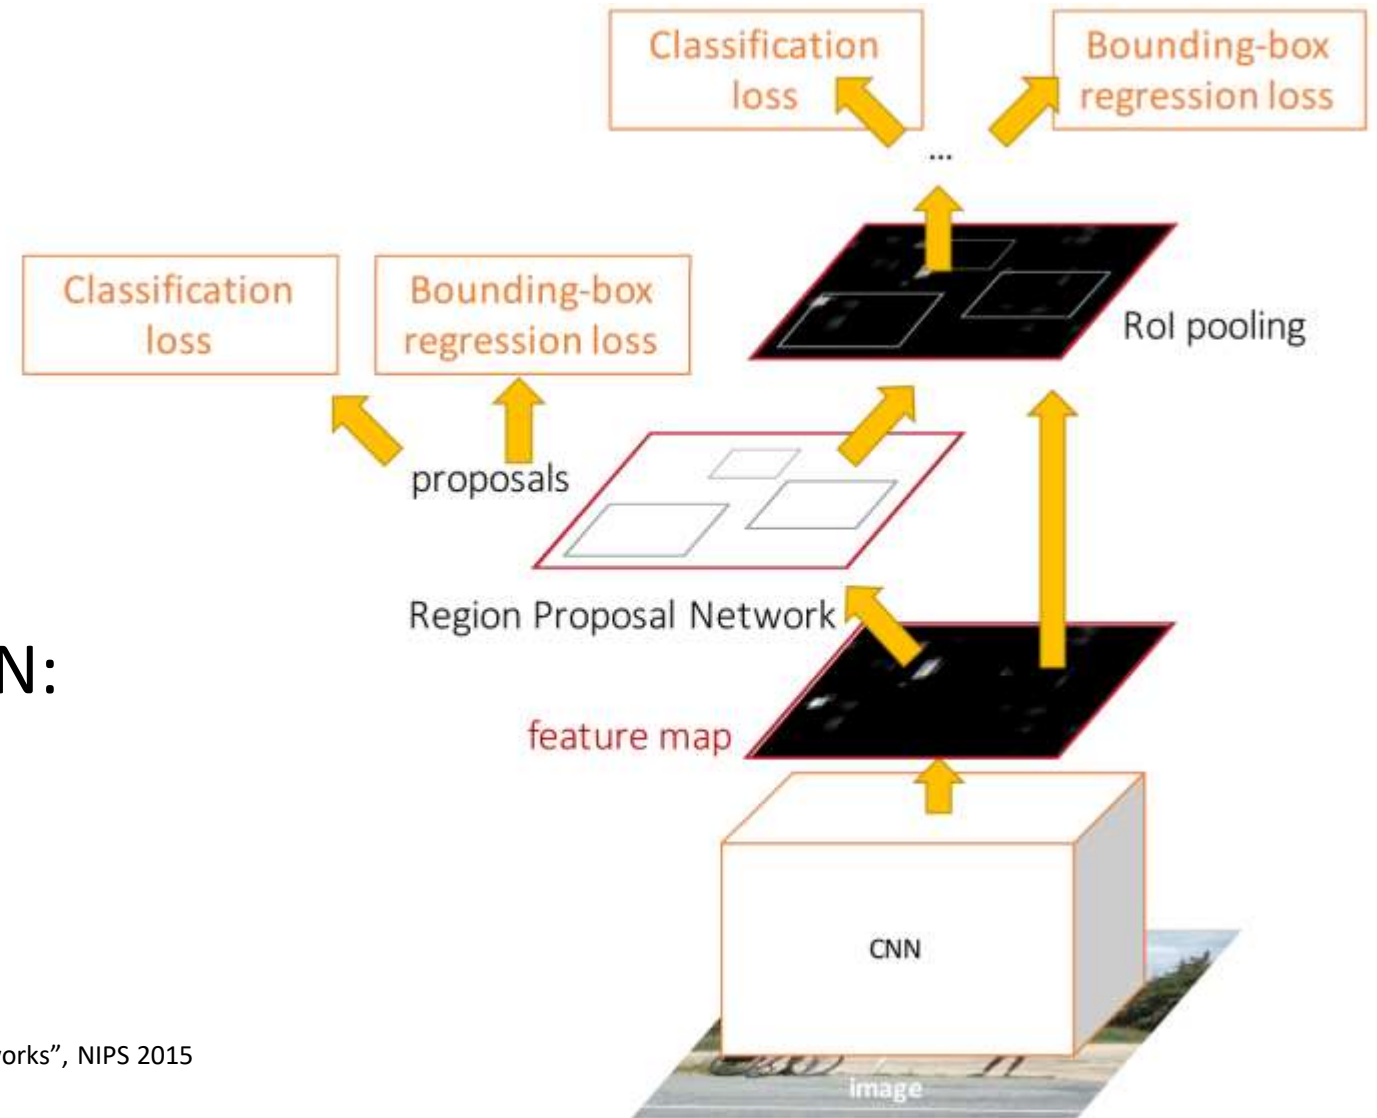
**Problem:** Runtime dominated by region proposals!

**Recall:** Region proposals computed by heuristic “Selective Search” algorithm on CPU -- let’s learn them with a CNN instead!

# Faster R-CNN: Learnable Region Proposals

Insert **Region Proposal Network (RPN)** to predict proposals from features

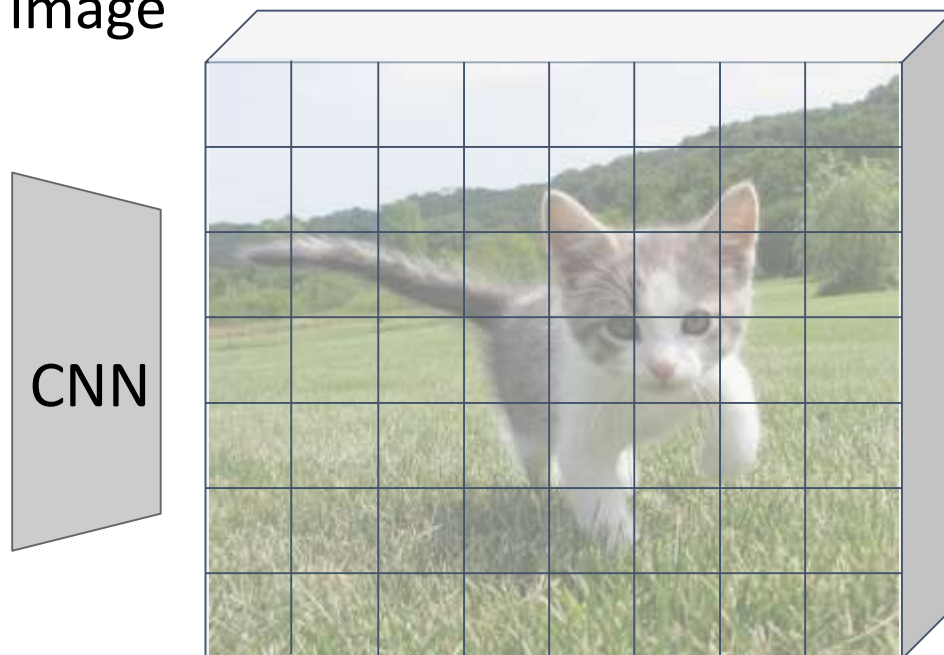
Otherwise same as Fast R-CNN:  
Crop features for each proposal, classify each one



Ren et al, "Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks", NIPS 2015  
Figure copyright 2015, Ross Girshick; reproduced with permission

# Region Proposal Network (RPN)

Run backbone CNN to get  
features aligned to input image



Input Image  
(e.g. 3 x 640 x 480)

Image features  
(e.g. 512 x 20 x 15)

# Region Proposal Network (RPN)

Run backbone CNN to get  
features aligned to input image



Input Image  
(e.g. 3 x 640 x 480)

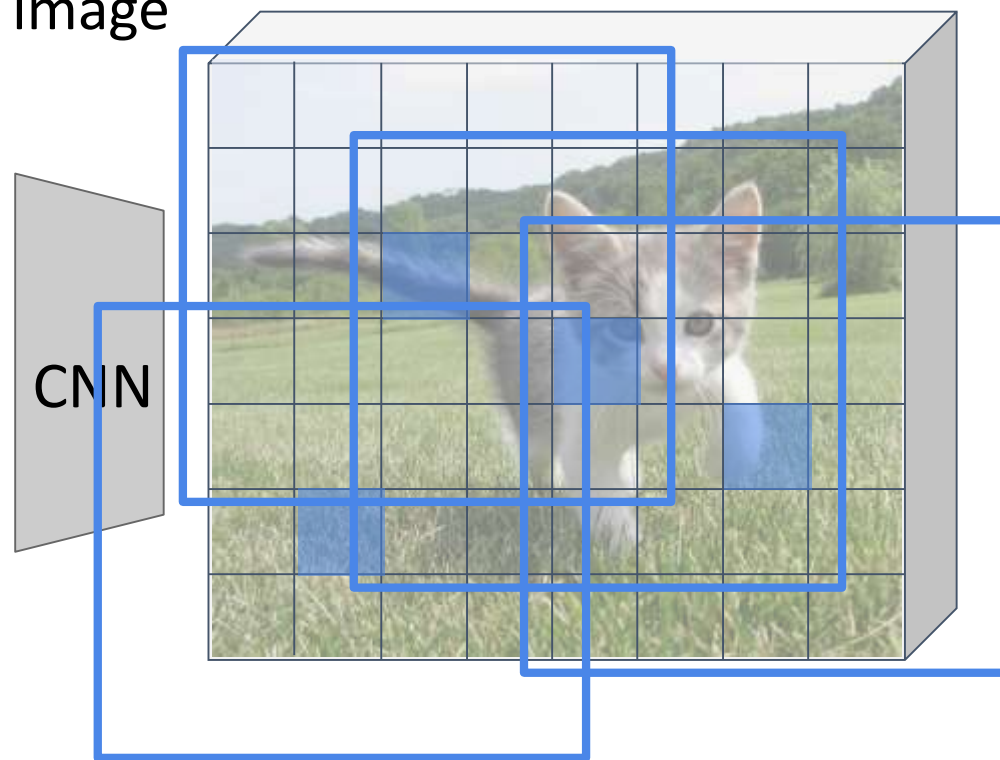


Image features  
(e.g. 512 x 20 x 15)

Imagine an **anchor box** of  
fixed size at each point in  
the feature map



# Region Proposal Network (RPN)

Run backbone CNN to get features aligned to input image



Input Image  
(e.g. 3 x 640 x 480)




Image features  
(e.g. 512 x 20 x 15)

Imagine an anchor box of fixed size at each point in the feature map

Anchor is an object?  
1 x 20 x 15

At each point, predict whether the corresponding anchor contains an object (per-cell logistic regression, predict scores with conv layer)



# Region Proposal Network (RPN)

Run backbone CNN to get features aligned to input image



Input Image  
(e.g.  $3 \times 640 \times 480$ )

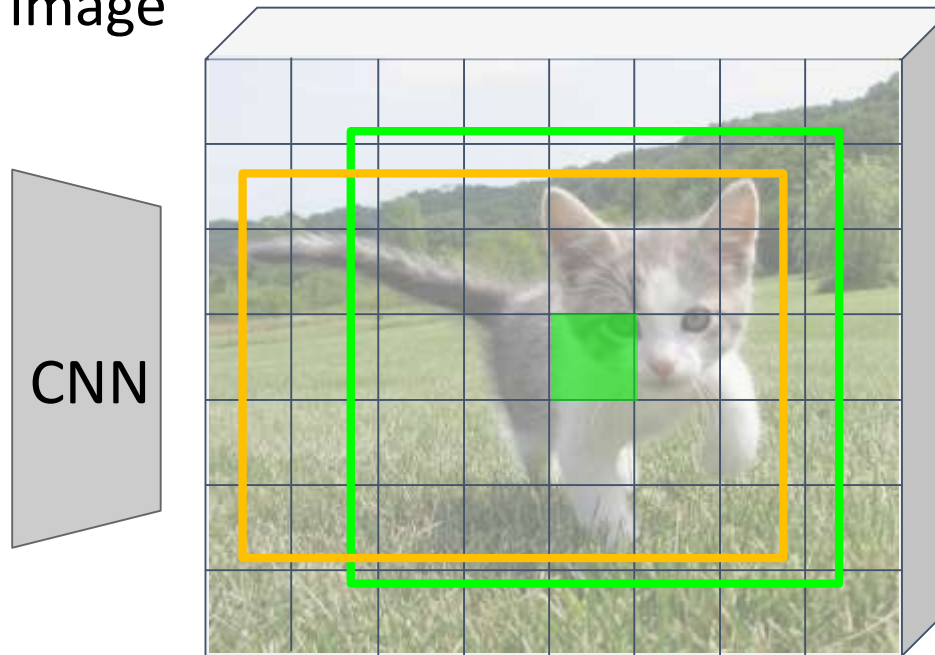


Image features  
(e.g.  $512 \times 20 \times 15$ )

Imagine an anchor box of fixed size at each point in the feature map



Anchor is an object?  
 $1 \times 20 \times 15$

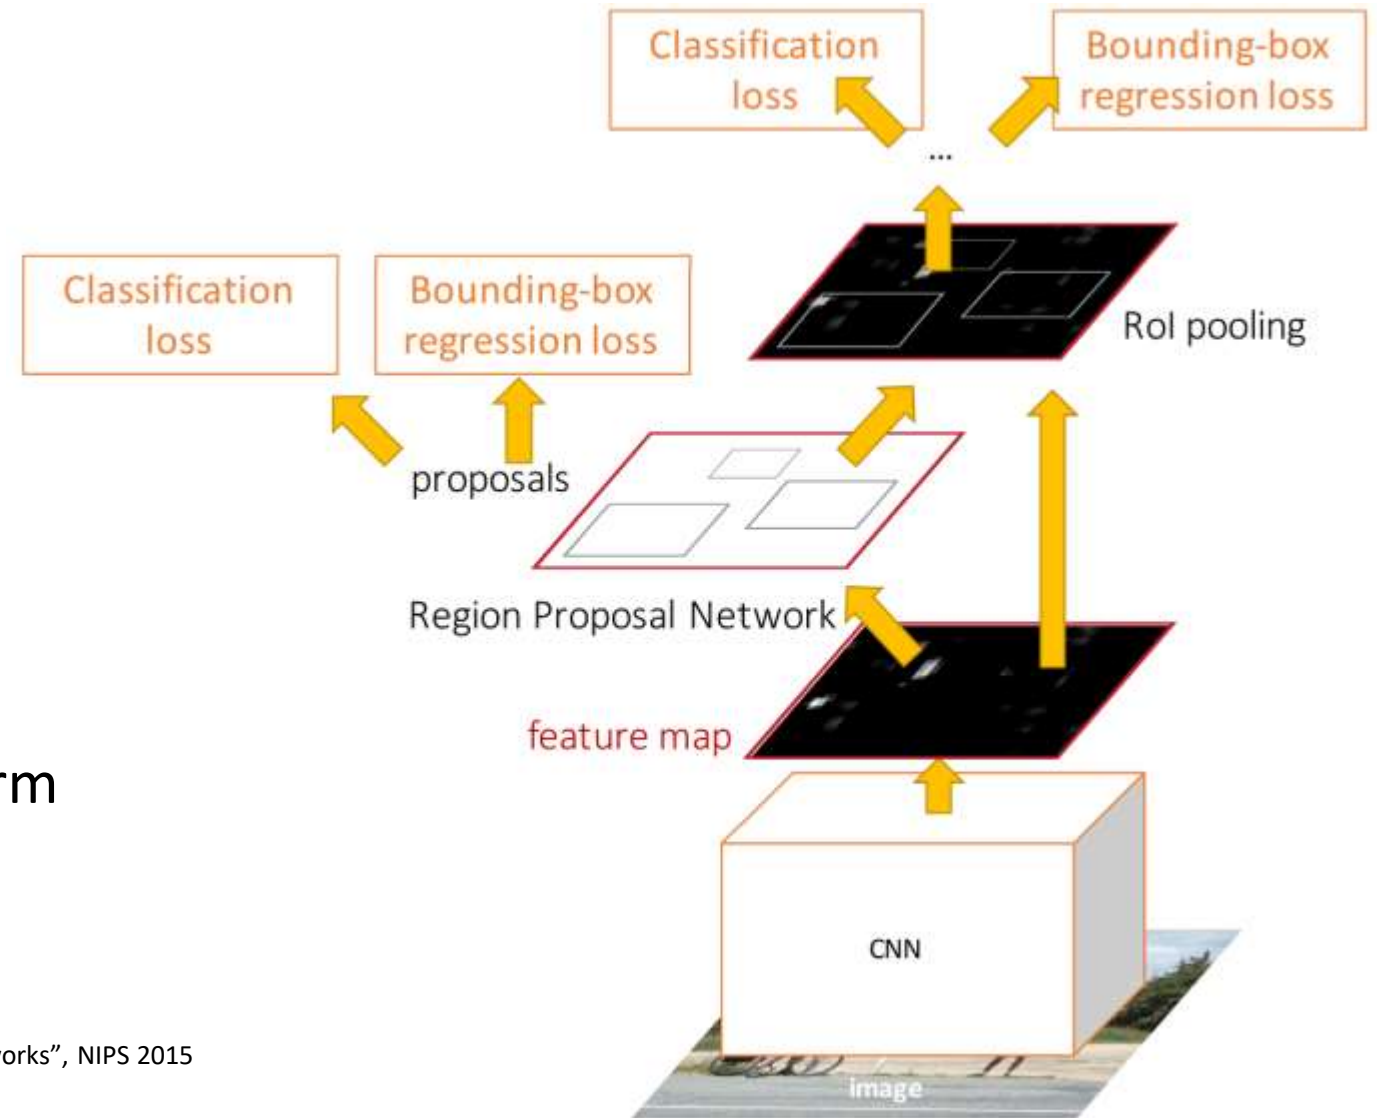
Box transforms  
 $4 \times 20 \times 15$

For positive boxes, also predict a box transform to regress from **anchor box** to **object box**

# Faster R-CNN: Learnable Region Proposals


Jointly train with 4 losses:

1. **RPN classification:** anchor box is object / not an object
2. **RPN regression:** predict transform from anchor box to proposal box
3. **Object classification:** classify proposals as background / object class
4. **Object regression:** predict transform from proposal box to object box



Ren et al, "Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks", NIPS 2015  
Figure copyright 2015, Ross Girshick; reproduced with permission

# Faster R-CNN: Learnable Region Proposals



# Faster R-CNN: Learnable Region Proposals

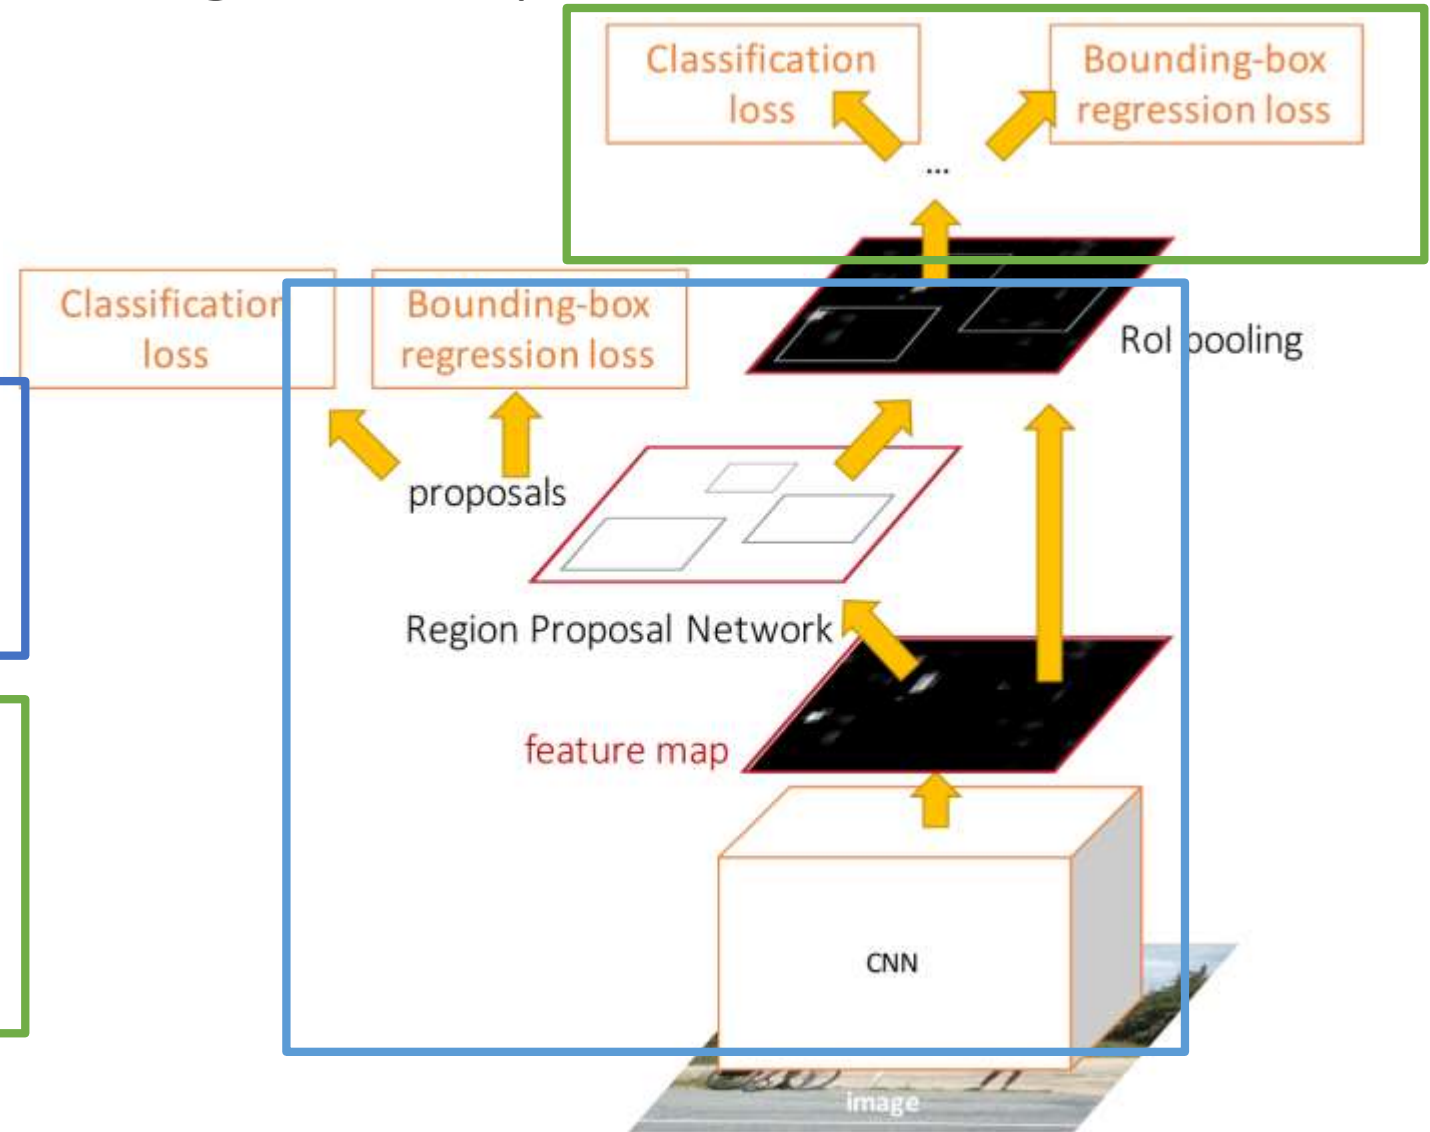
Faster R-CNN is a  
**Two-stage object detector**

First stage: Run once per image

- Backbone network
- Region proposal network

Second stage: Run once per region

- Crop features: ROI pool / align
- Predict object class
- Prediction bbox offset



# Faster R-CNN: Learnable Region Proposals

Question: **Do we really need the second stage?**

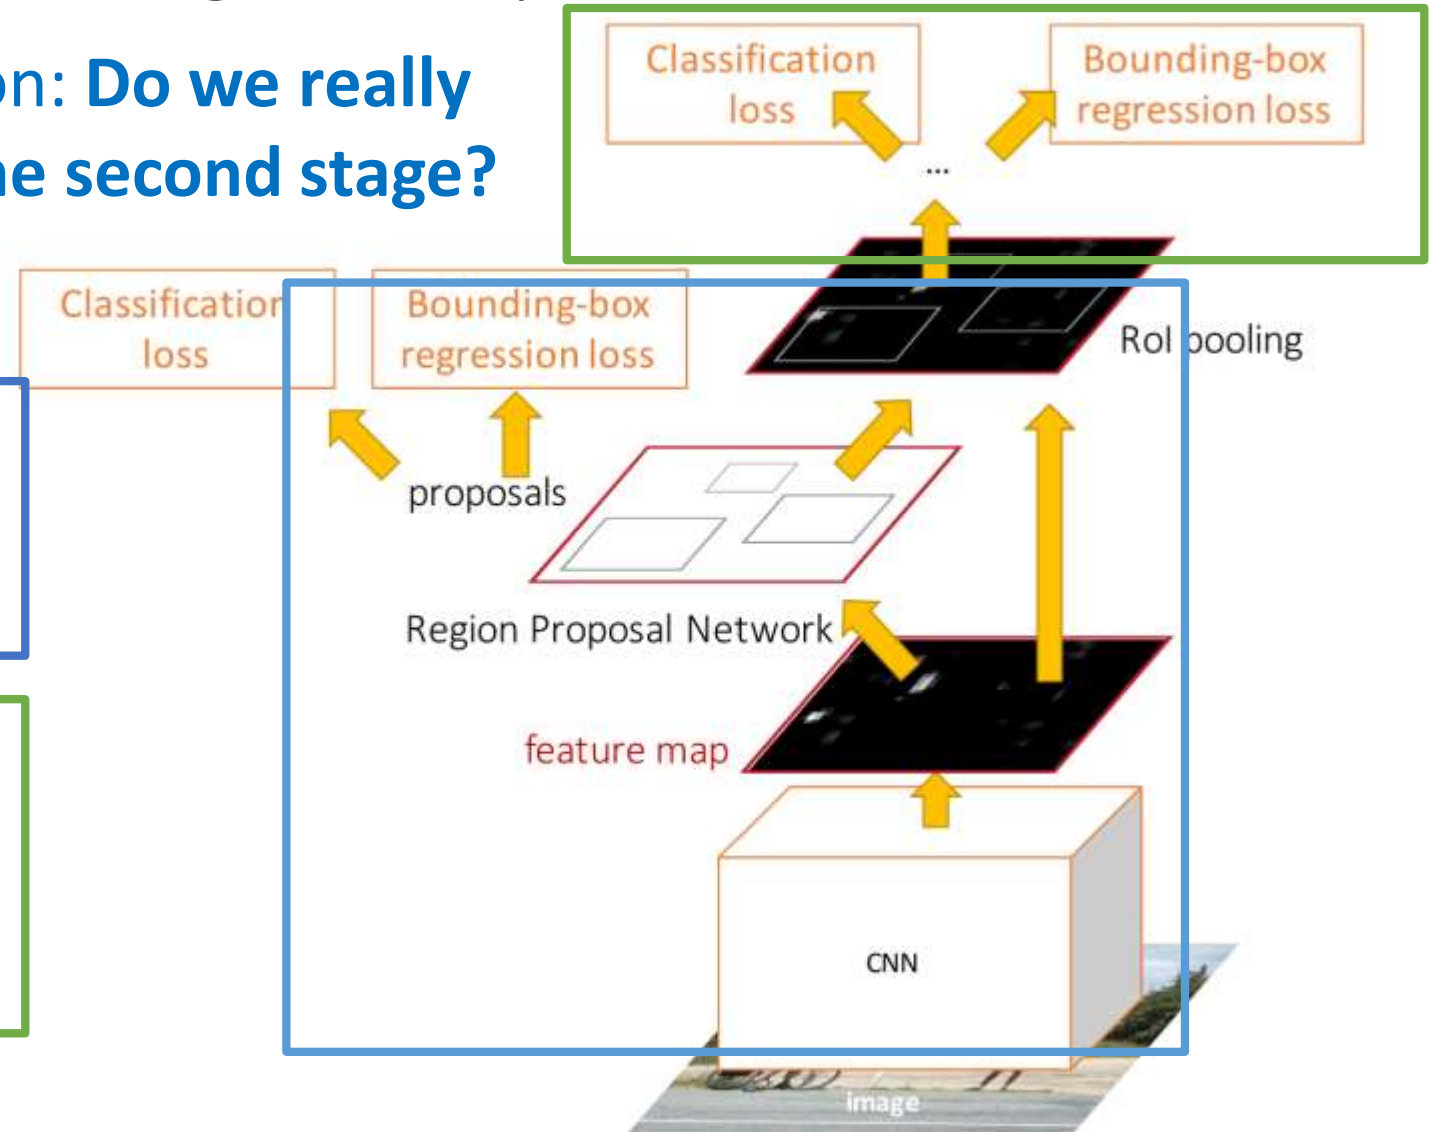
Faster R-CNN is a  
**Two-stage object detector**

First stage: Run once per image

- Backbone network
- Region proposal network

Second stage: Run once per region

- Crop features: RoI pool / align
- Predict object class
- Prediction bbox offset





# Single-Stage Object Detection

Run backbone CNN to get features aligned to input image



Input Image  
(e.g. 3 x 640 x 480)





Image features  
(e.g. 512 x 20 x 15)

**RPN:** Classify each anchor as object / not object  
**Single-Stage Detector:** Classify each object as one of C categories (or background)



Remember: K anchors at each position in image feature map

# Single-Stage Object Detection

Run backbone CNN to get features aligned to input image



Input Image  
(e.g. 3 x 640 x 480)

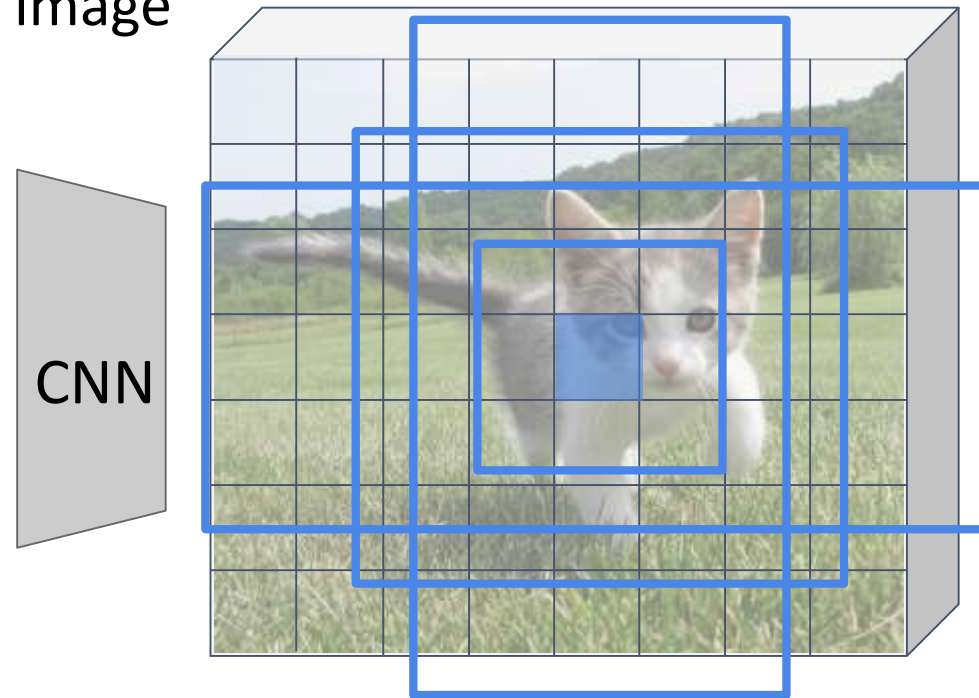



Image features  
(e.g. 512 x 20 x 15)


**RPN:** Classify each anchor as object / not object  
**Single-Stage Detector:** Classify each object as one of C categories (or background)



Sometimes use **category-specific regression**: Predict different box transforms for each category


Redmon et al, "You Only Look Once: Unified, Real-Time Object Detection", CVPR 2016  
Liu et al, "SSD: Single-Shot MultiBox Detector", ECCV 2016  
Lin et al, "Focal Loss for Dense Object Detection", ICCV 2017

# Object Detection: Lots of variables!



Huang et al, "Speed/accuracy trade-offs for modern convolutional object detectors", CVPR 2017

# Object Detection: Lots of variables!




## Takeaways:

- Two stage method (Faster R-CNN) get the best accuracy, but are slower

Huang et al, "Speed/accuracy trade-offs for modern convolutional object detectors", CVPR 2017



# Object Detection: Lots of variables!




## Takeaways:

- Two stage method (Faster R-CNN) get the best accuracy, but are slower
- Single-stage methods (SSD) are much faster, but don't perform as well

Huang et al, "Speed/accuracy trade-offs for modern convolutional object detectors", CVPR 2017



# Object Detection: Lots of variables!




## Takeaways:

- Two stage method (Faster R-CNN) get the best accuracy, but are slower
- Single-stage methods (SSD) are much faster, but don't perform as well
- Bigger backbones improve performance, but are slower

Huang et al, "Speed/accuracy trade-offs for modern convolutional object detectors", CVPR 2017

# Object Detection: Lots of variables!

These results are a few years old ... since then GPUs have gotten faster, and we've improved performance with many tricks:




Huang et al, "Speed/accuracy trade-offs for modern convolutional object detectors", CVPR 2017

Wu et al, Detectron2, GitHub 2019

# Object Detection: Lots of variables!

Faster R-CNN  
w/ResNet-101-FPN,  
longer training  
(63ms, 42.0 mAP)




These results are a few years old ... since then GPUs have gotten faster, and we've improved performance with many tricks:

- Train longer!
- Multiscale backbone: Feature Pyramid Networks

Huang et al, "Speed/accuracy trade-offs for modern convolutional object detectors", CVPR 2017

Wu et al, Detectron2, GitHub 2019

# Object Detection: Lots of variables!




These results are a few years old ... since then GPUs have gotten faster, and we've improved performance with many tricks:

- Train longer!
- Multiscale backbone: Feature Pyramid Networks
- Better backbone: ResNeXt

Huang et al, "Speed/accuracy trade-offs for modern convolutional object detectors", CVPR 2017

Wu et al, Detectron2, GitHub 2019

# Object Detection: Lots of variables!



These results are a few years old ... since then GPUs have gotten faster, and we've improved performance with many tricks:


- Train longer!
- Multiscale backbone: Feature Pyramid Networks
- Better backbone: ResNeXt
- Single-Stage methods have improved

Huang et al, "Speed/accuracy trade-offs for modern convolutional object detectors", CVPR 2017

Wu et al, Detectron2, GitHub 2019



# Object Detection: Lots of variables!




These results are a few years old ... since then GPUs have gotten faster, and we've improved performance with many tricks:

- Train longer!
- Multiscale backbone: Feature Pyramid Networks
- Better backbone: ResNeXt
- Single-Stage methods have improved
- Very big models work better


Huang et al, "Speed/accuracy trade-offs for modern convolutional object detectors", CVPR 2017

Wu et al, Detectron2, GitHub 2019

# Object Detection: Lots of variables!



# Object Detection: Lots of variables!



These results are a few years old ... since then GPUs have gotten faster, and we've improved performance with many tricks:

- Train longer!
- Multiscale backbone: Feature Pyramid Networks
- Better backbone: ResNeXt
- Single-Stage methods have improved
- Very big models work better
- Test-time augmentation pushes numbers up

Huang et al, "Speed/accuracy trade-offs for modern convolutional object detectors", CVPR 2017

Wu et al, Detectron2, GitHub 2019

# Object Detection: Open-Source Code

## **TensorFlow Detection API:**

[https://github.com/tensorflow/models/tree/master/research/object\\_detection](https://github.com/tensorflow/models/tree/master/research/object_detection)

Faster R-CNN, SSD, RFCN, Mask R-CNN

## **Detectron2 (PyTorch):**

<https://github.com/facebookresearch/detectron2>

Fast / Faster / Mask R-CNN, RetinaNet

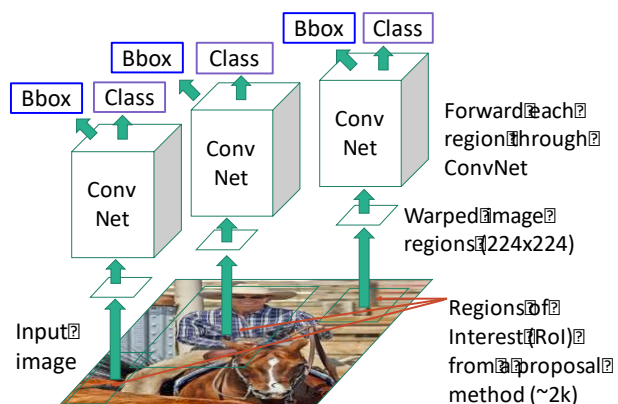
## **YOLOv5 (Ultralytics – PyTorch):**

<https://github.com/ultralytics/yolov5>

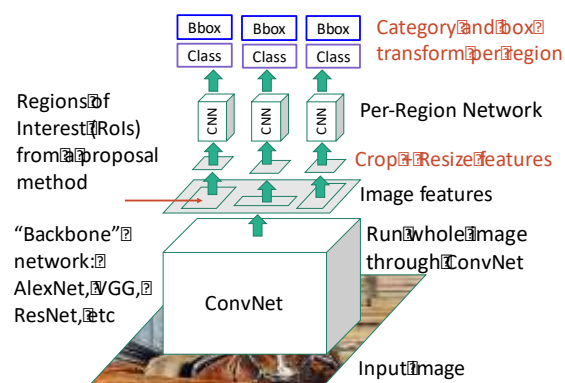
Single-stage

# Summary


**“Slow” R-CNN:** Run CNN independently for each region



**Fast R-CNN:** Apply differentiable cropping to shared image features



**Faster R-CNN:** Compute proposals with CNN



**Single-Stage:** Fully convolutional detector





Test 2  
20 Dec, Monday  
900-1100 am  
Online via Teams